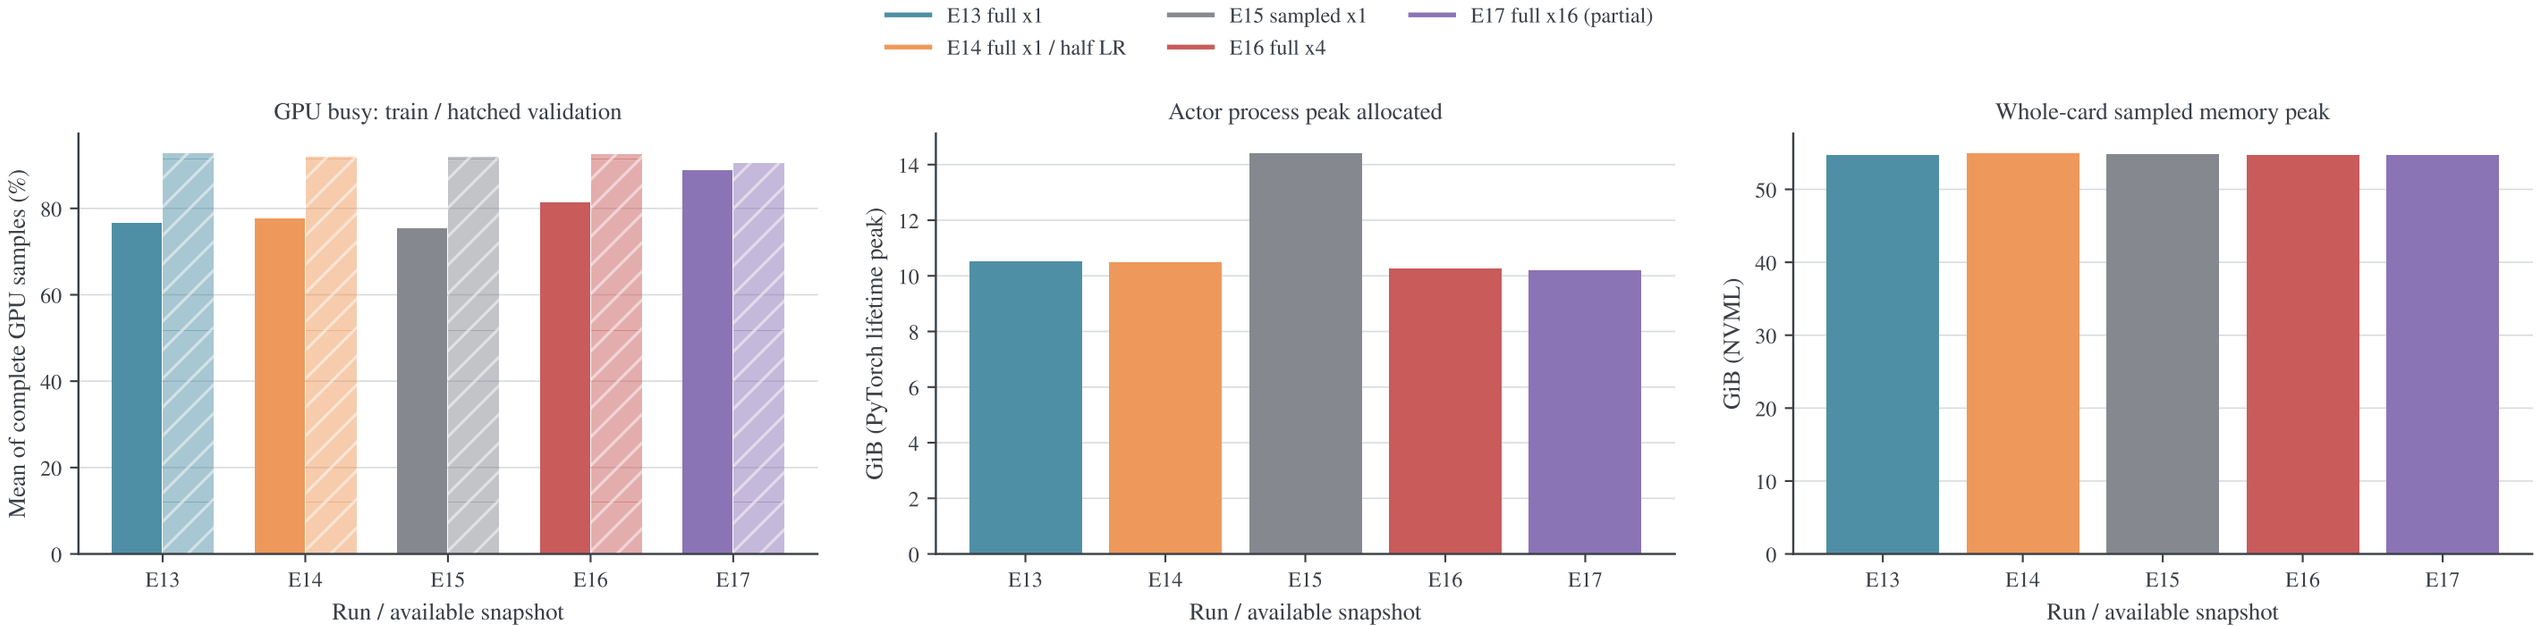
E13 full x1
E14 full x1 / half LR
E15 sampled x1
E16 full x4
E17 full x16 (partial)
GPU busy: train / hatched validation
Mean of complete GPU samples (%)
0
20
40
60
80
E13
E14
E15
E16
E17
Run / available snapshot
Actor process peak allocated
GiB (PyTorch lifetime peak)
0
2
4
6
8
10
12
14
E13
E14
E15
E16
E17
Run / available snapshot
Whole-card sampled memory peak
GiB (NVML)
0
10
20
30
40
50
E13
E14
E15
E16
E17
Run / available snapshot
Approximate phases with 10-second edge trimming; periodic samples can miss peaks; GPU busy is not MFU