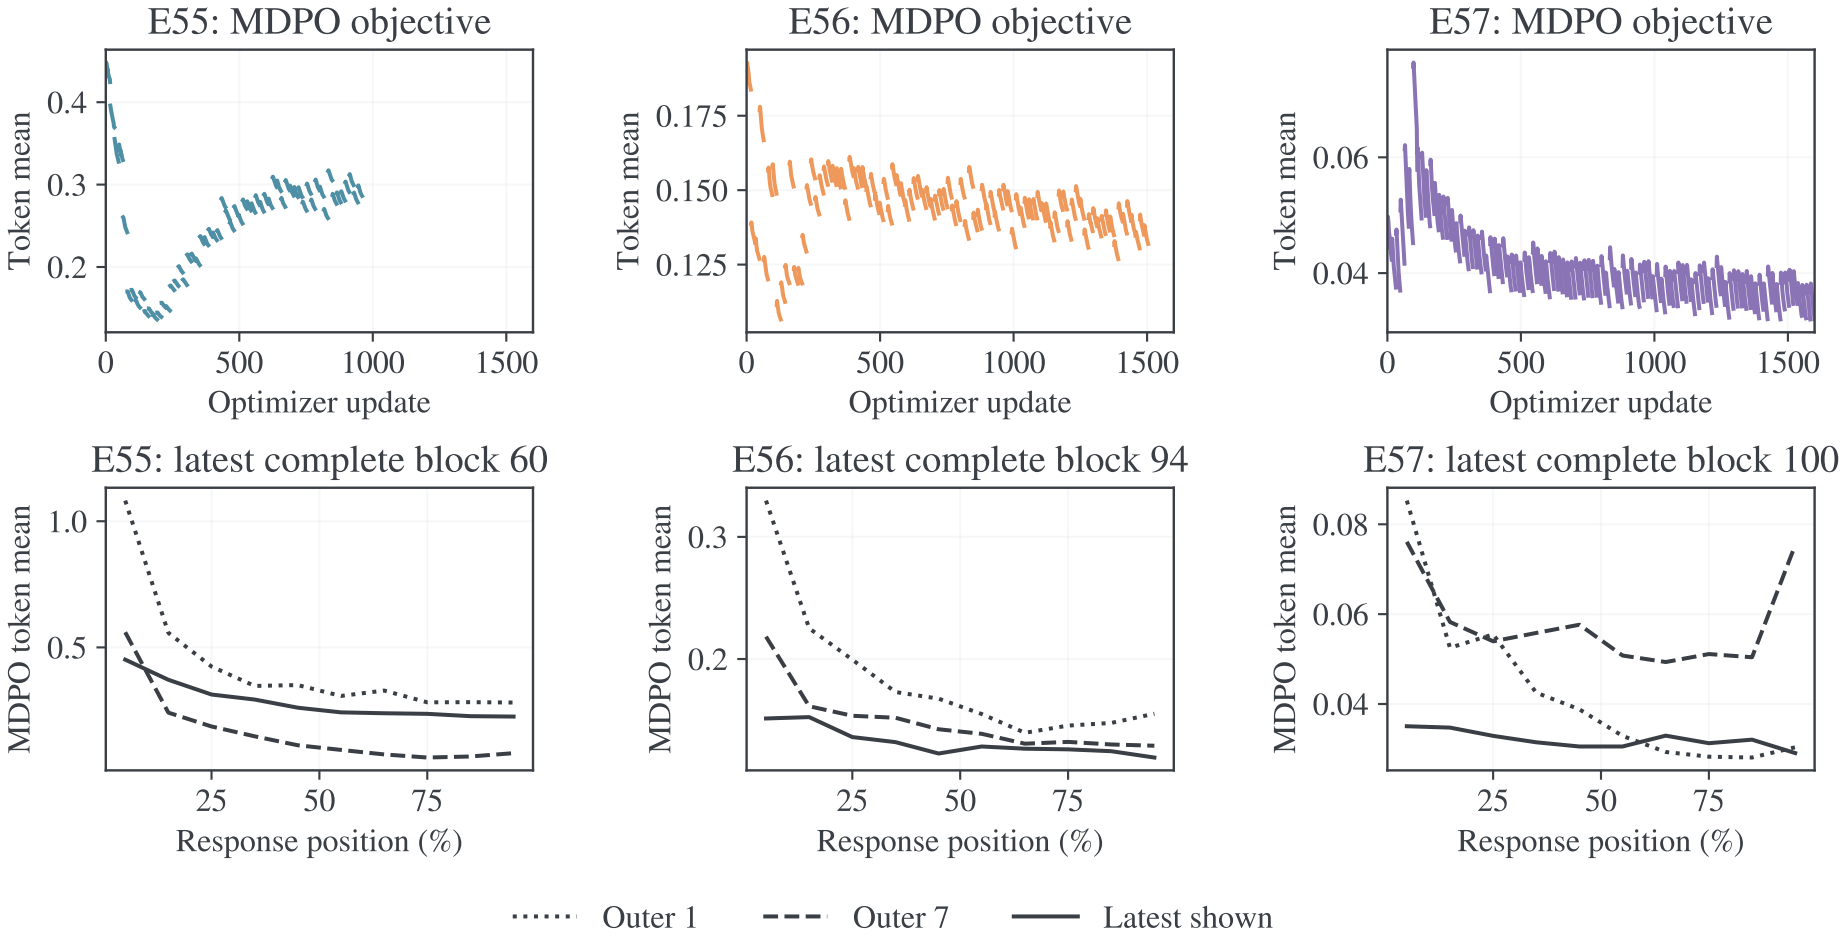
Qwen3-1.7B-Base <- Qwen3-30B-A3B-Instruct-2507
Optimization and MDPO position profiles
E55: MDPO objective
E56: MDPO objective
E57: MDPO objective
Token mean
Optimizer update
E55: latest complete block 60
E56: latest complete block 94
E57: latest complete block 100
MDPO token mean
Response position (%)
Outer 1
Outer 7
Latest shown
Each segment is one fixed batch; targets differ across kappa, so raw losses are not ranks.
Bottom: before each selected outer's 16th update; these are not teacher-KL position curves.