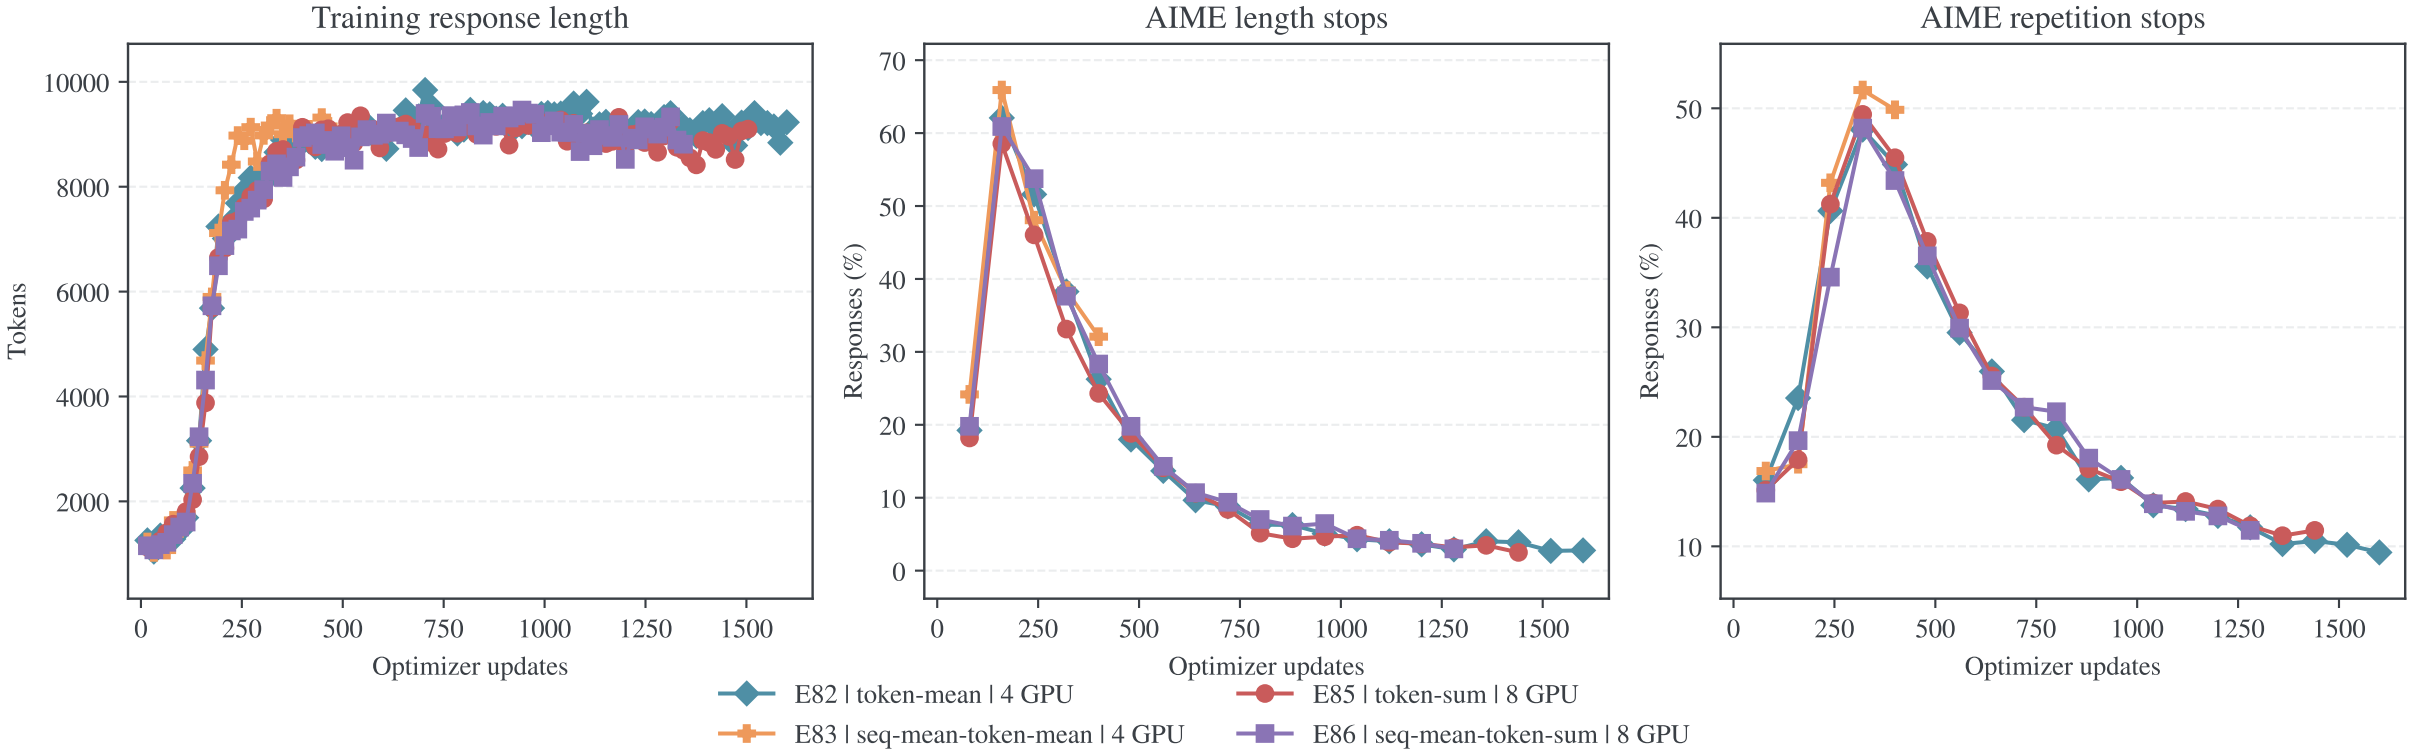
Training response length
Tokens
Optimizer updates
AIME length stops
Responses (%)
Optimizer updates
AIME repetition stops
Responses (%)
Optimizer updates
E82 | token-mean | 4 GPU
E83 | seq-mean-token-mean | 4 GPU
E85 | token-sum | 8 GPU
E86 | seq-mean-token-sum | 8 GPU
Training length and AIME validation stop fractions; complete observed nodes only