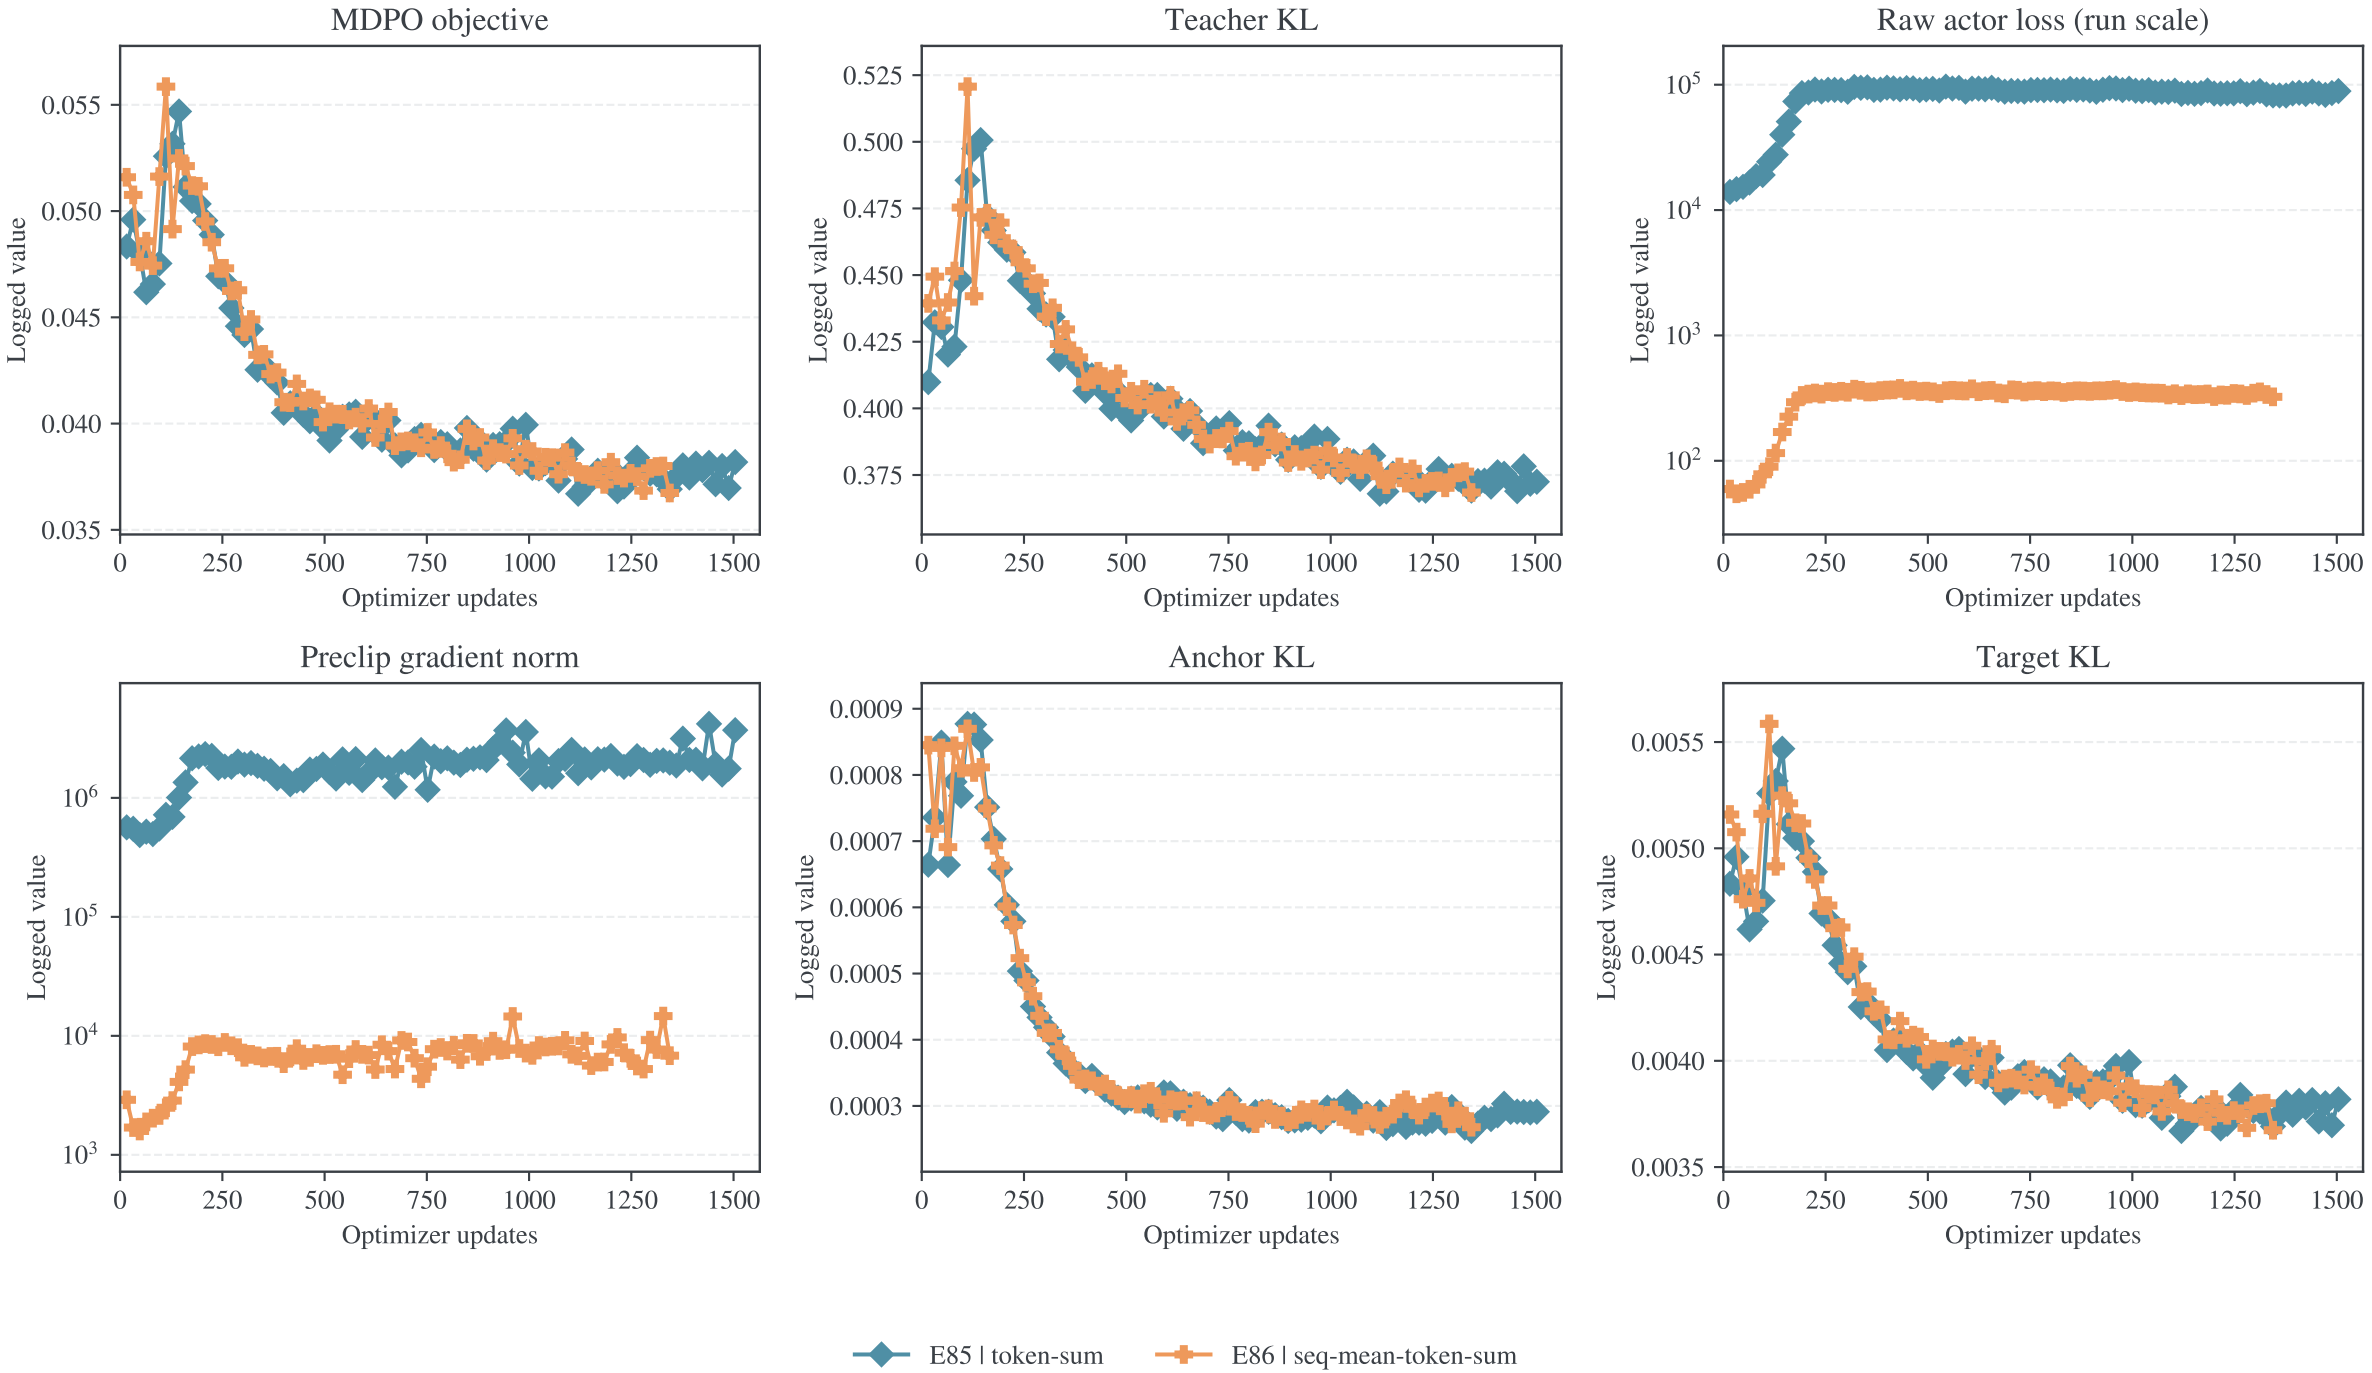
MDPO objective
Teacher KL
Raw actor loss (run scale)
Preclip gradient norm
Anchor KL
Target KL
Logged value
Optimizer updates
E85 | token-sum
E86 | seq-mean-token-sum
Objective and KL are token means; raw actor loss uses each run's aggregation and is not directly rankable