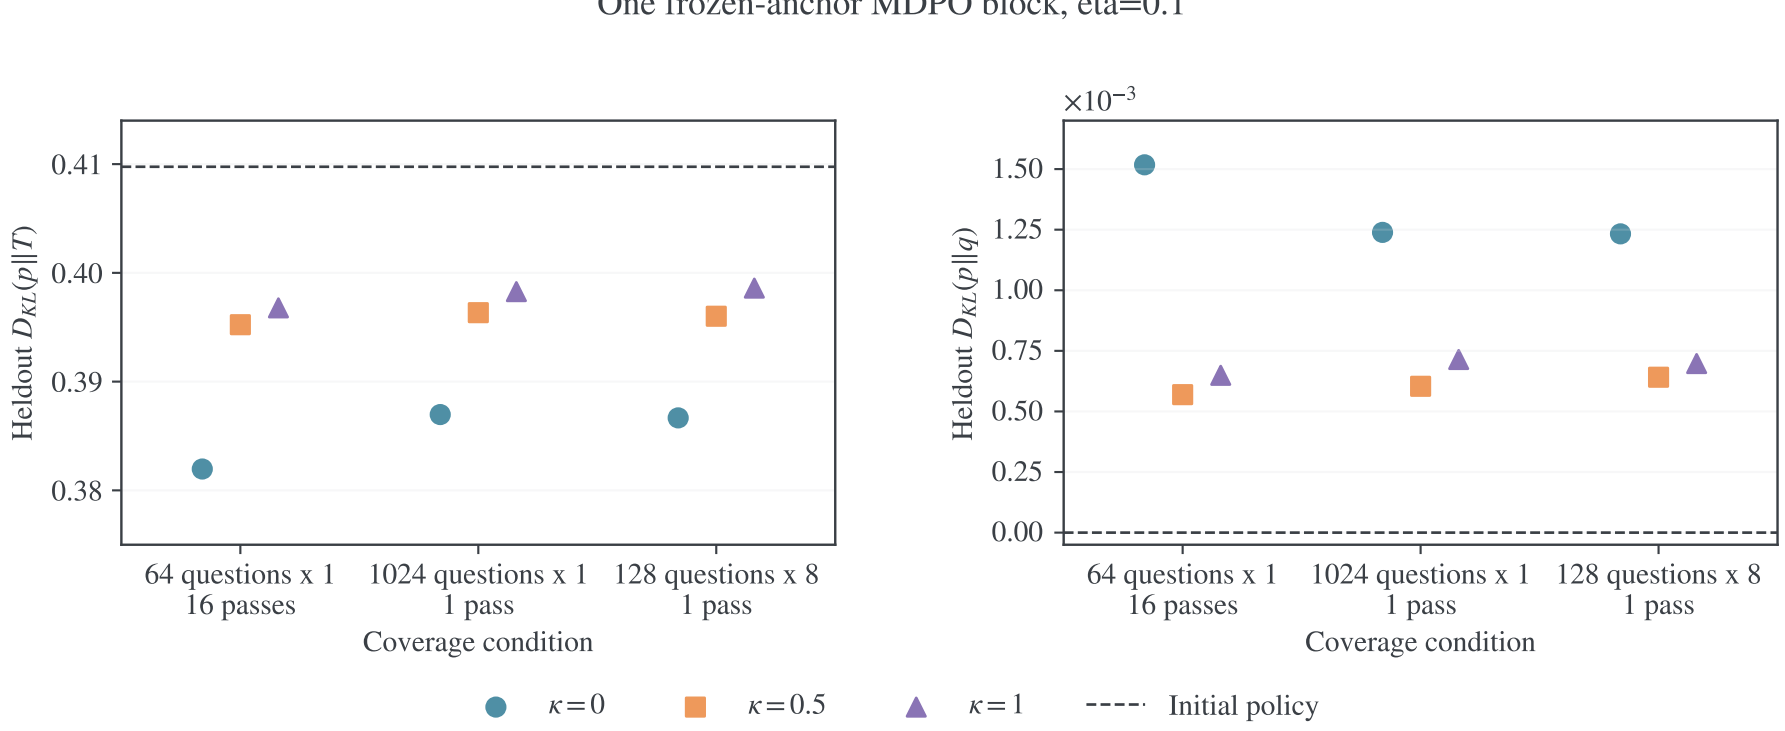
Qwen3-1.7B-Base <- Qwen3-30B-A3B-Instruct-2507
One frozen-anchor MDPO block, eta=0.1
Heldout $D_{KL}(p\|T)$
0.41
0.40
0.39
0.38
Heldout $D_{KL}(p\|q)$
$\times 10^{-3}$
1.50
1.25
1.00
0.75
0.50
0.25
0.00
64 questions x 1
16 passes
1024 questions x 1
1 pass
128 questions x 8
1 pass
Coverage condition
$\kappa = 0$
$\kappa = 0.5$
$\kappa = 1$
Initial policy
128 heldout Base trajectories; 16 updates x 64 responses; one bank/seed.
Fixed-prefix KL measurements; independent task evaluations are still required.