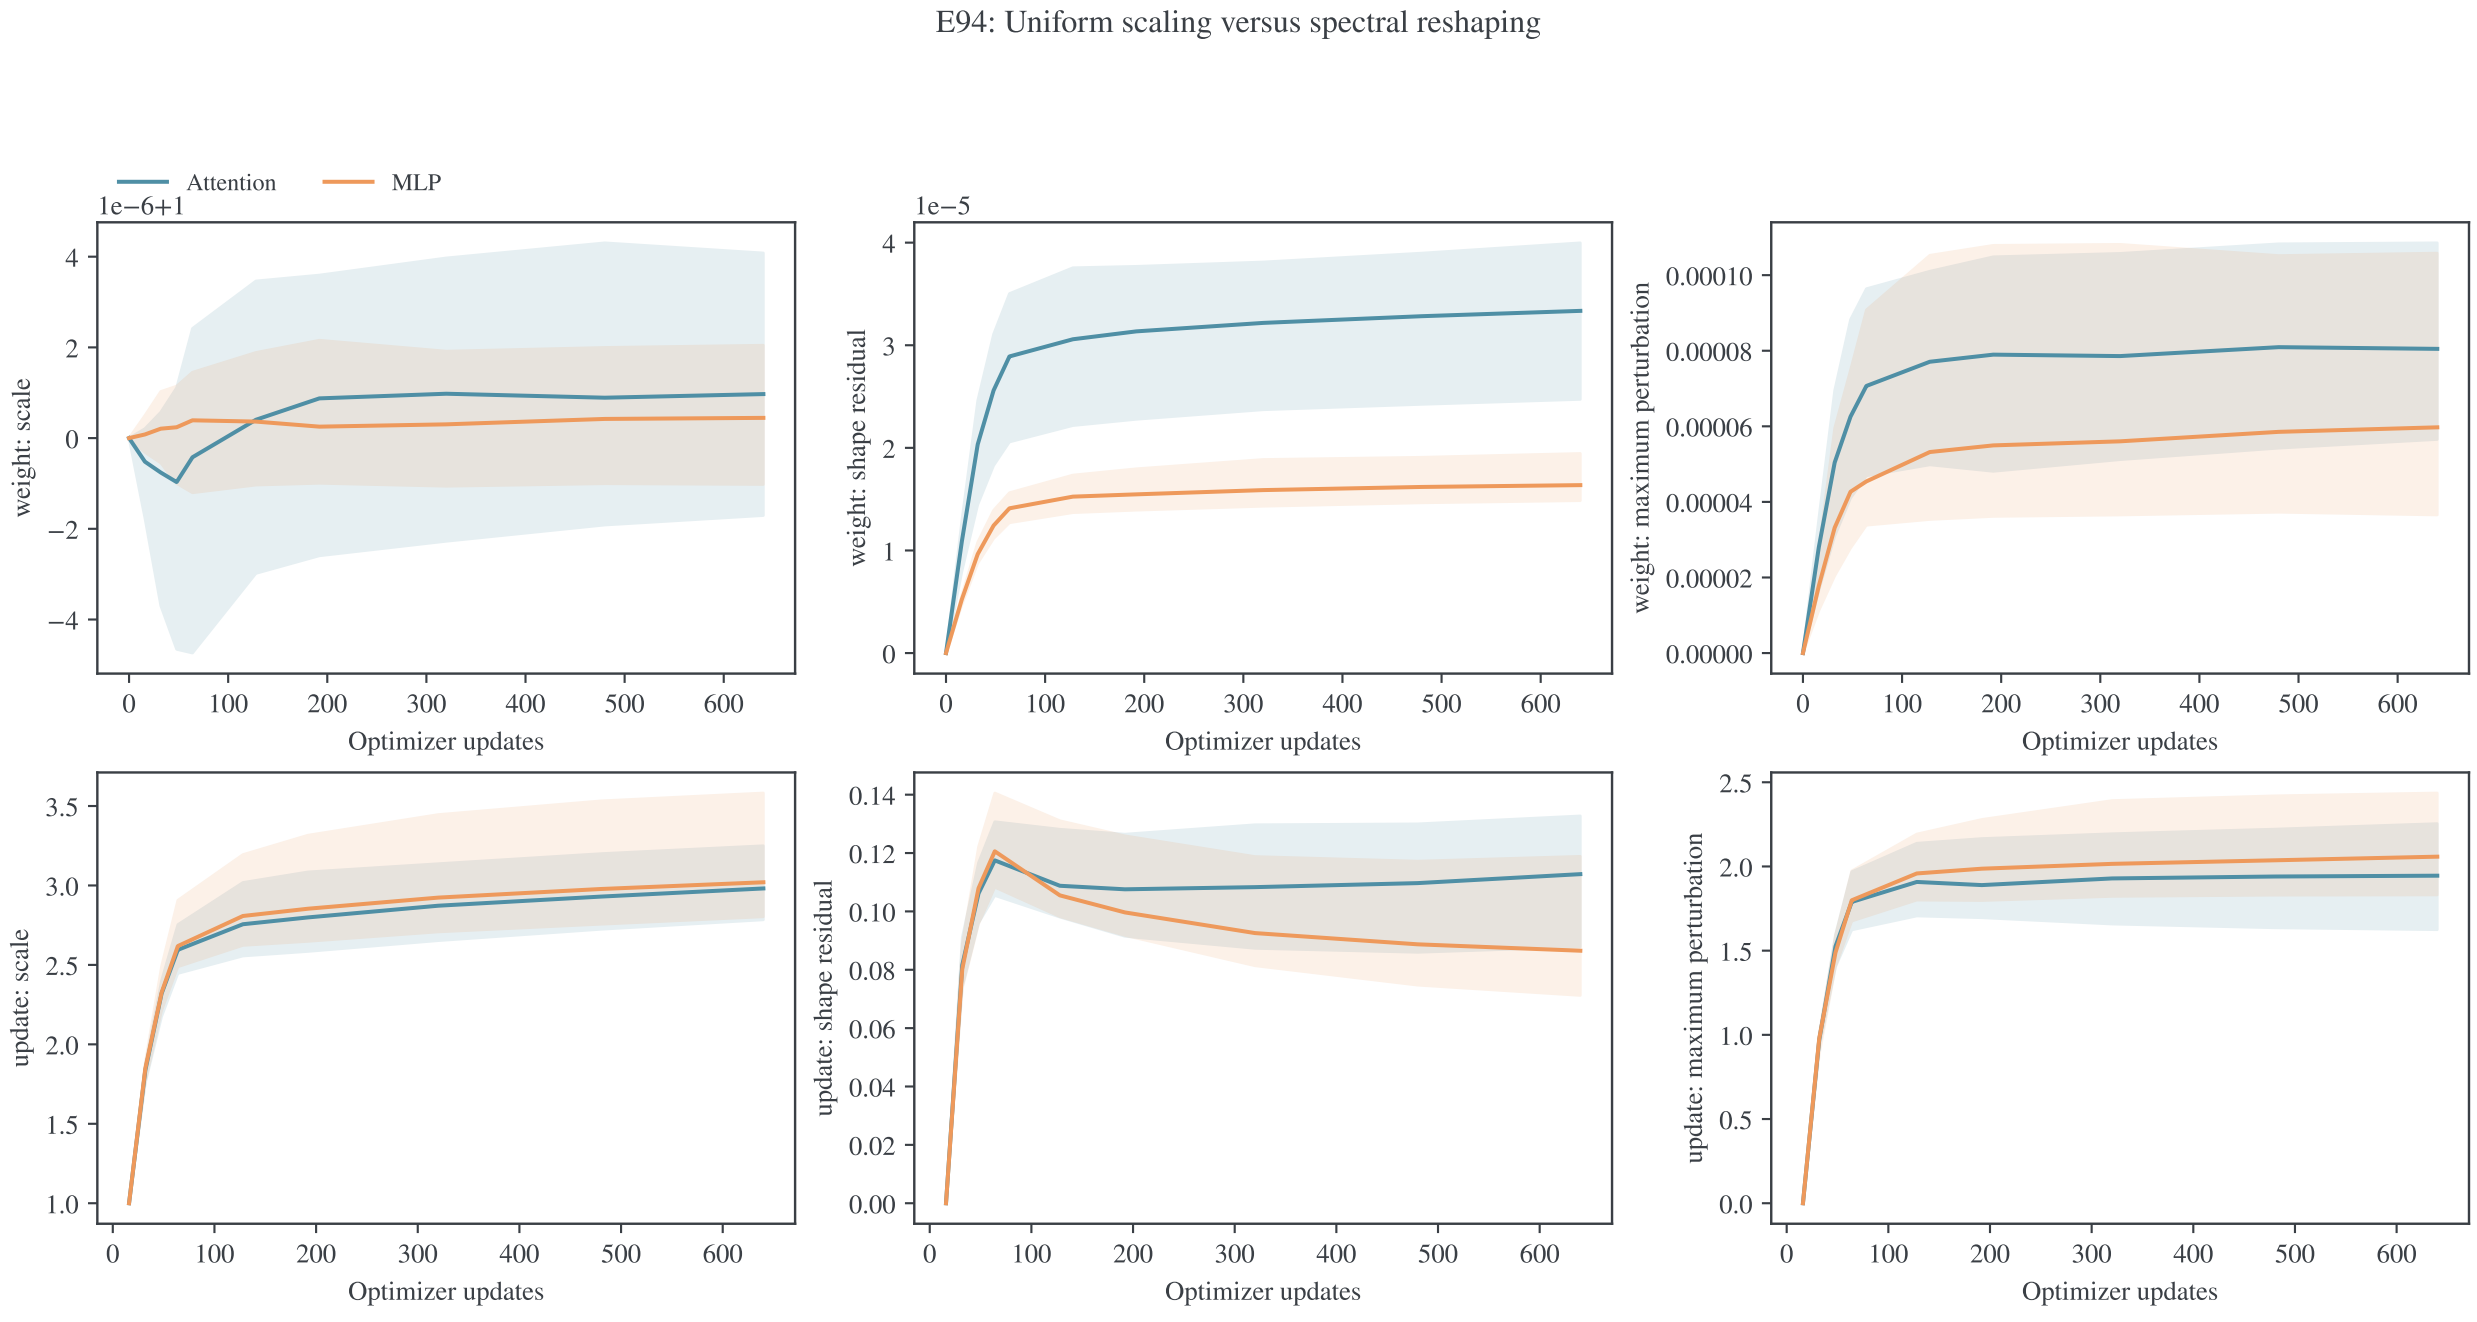
E94: Uniform scaling versus spectral reshaping
Attention
MLP
1e−6+1
weight: scale
1e−5
weight: shape residual
weight: maximum perturbation
update: scale
update: shape residual
update: maximum perturbation
Optimizer updates
196/196 matrices; 28 full layers. Weight reference: step 0. Cumulative-update reference: step 16. Matrix medians and IQR.