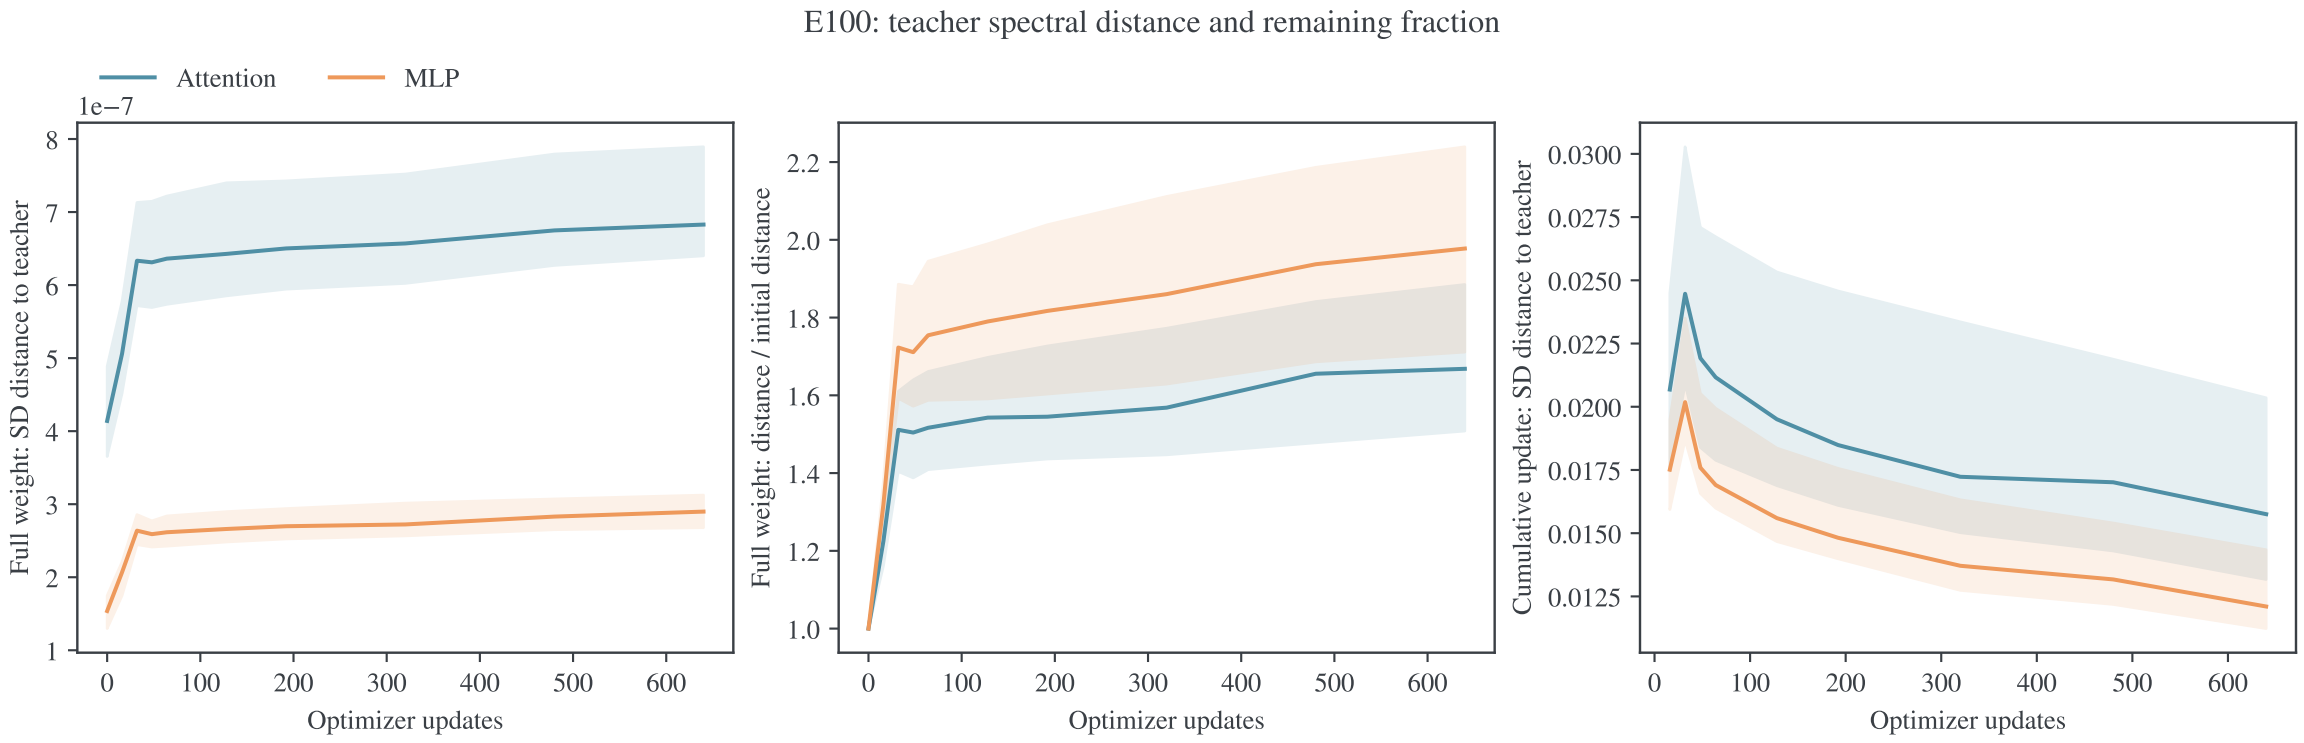
E100: teacher spectral distance and remaining fraction
Attention
MLP
1e−7
Full weight: SD distance to teacher
Optimizer updates
Full weight: distance / initial distance
Optimizer updates
Cumulative update: SD distance to teacher
Optimizer updates
Median and IQR; 196/196 matrices; 28 full layers. Relative distance excludes near-zero baselines.