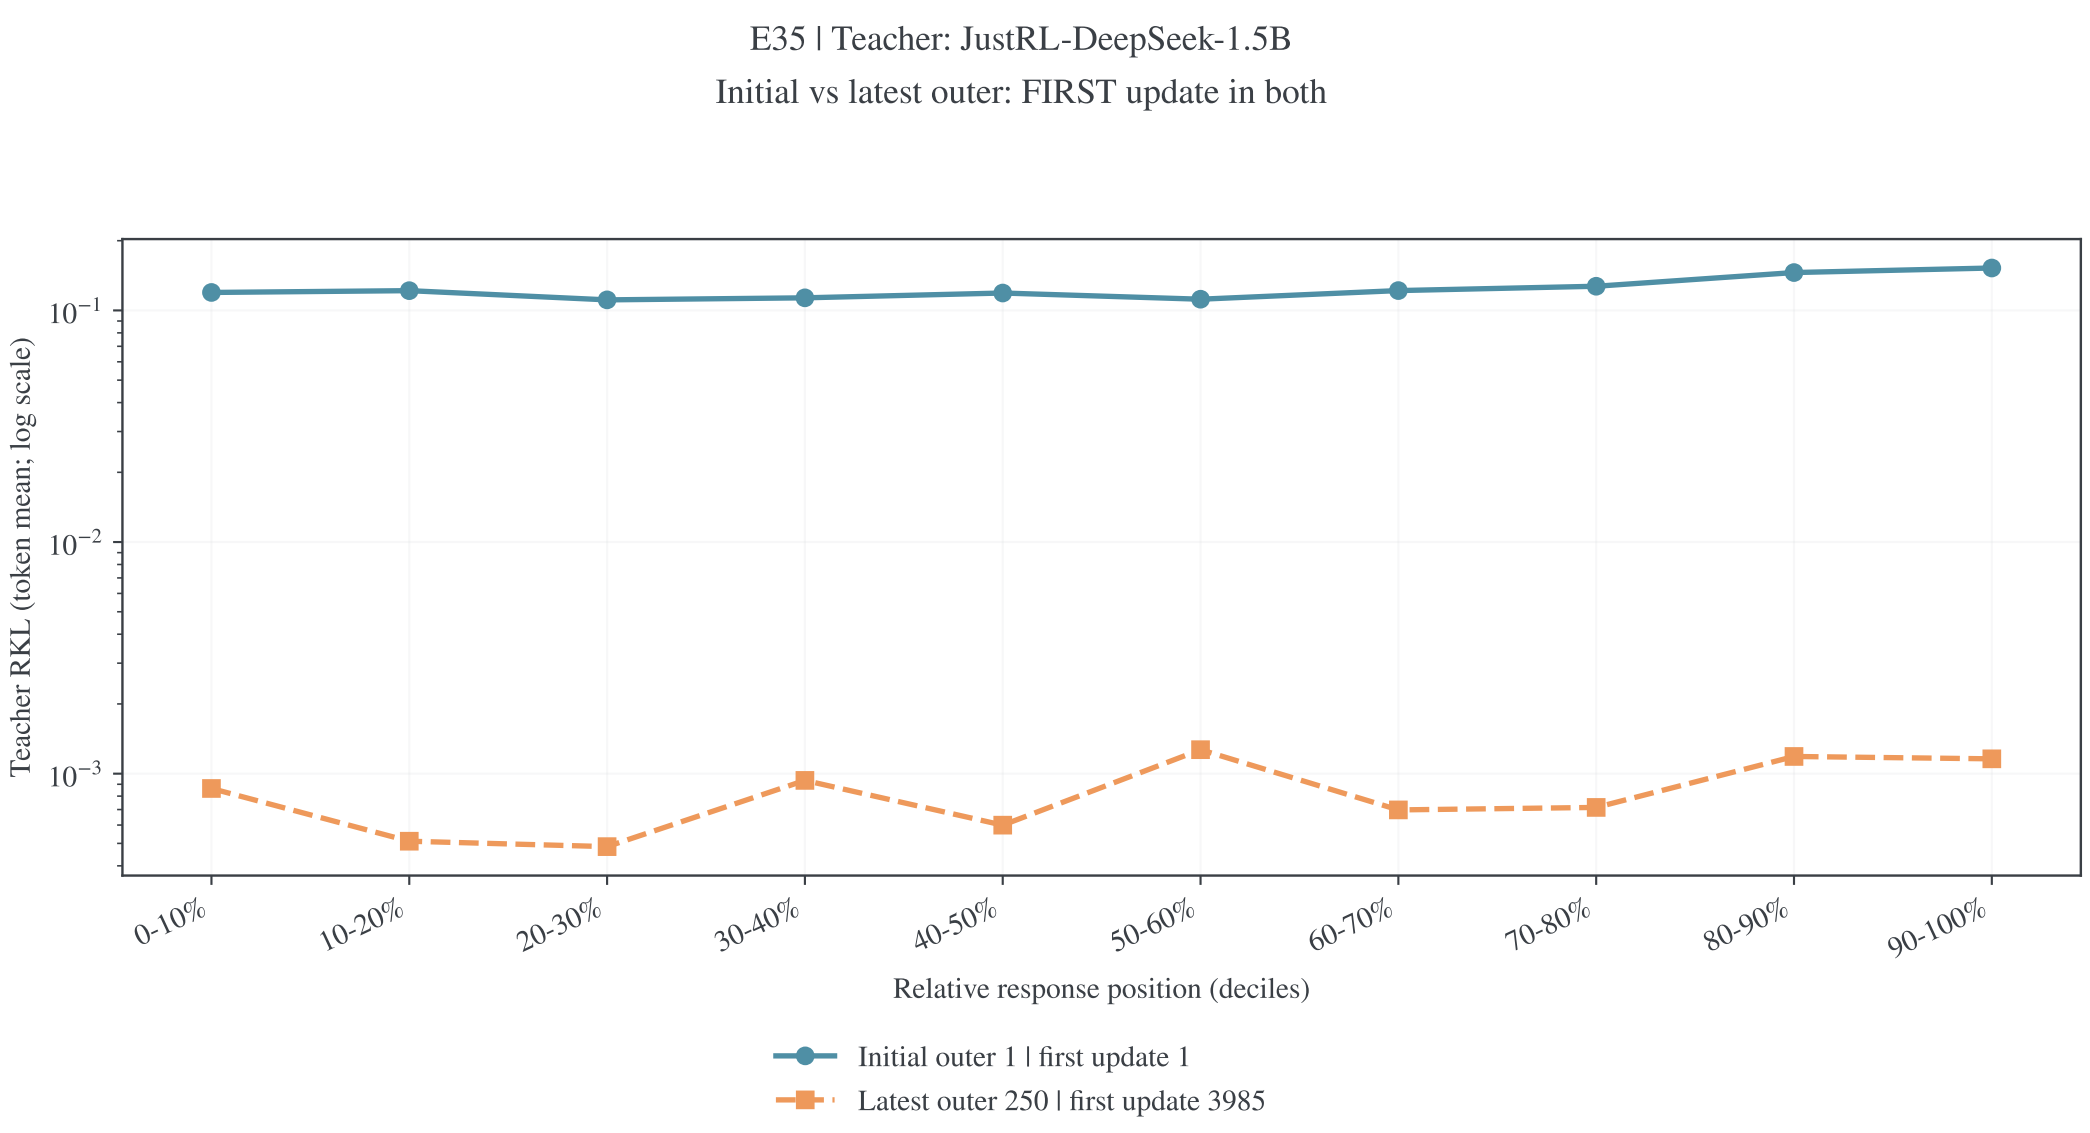
E35 | Teacher: JustRL-DeepSeek-1.5B
Initial vs latest outer: FIRST update in both
Teacher RKL (token mean; log scale)
$10^{-1}$
$10^{-2}$
$10^{-3}$
0-10%
10-20%
20-30%
30-40%
40-50%
50-60%
60-70%
70-80%
80-90%
90-100%
Relative response position (deciles)
Initial outer 1 | first update 1
Latest outer 250 | first update 3985
Both observations precede the first optimizer operation on their own rollout batch; no endpoint averaging.