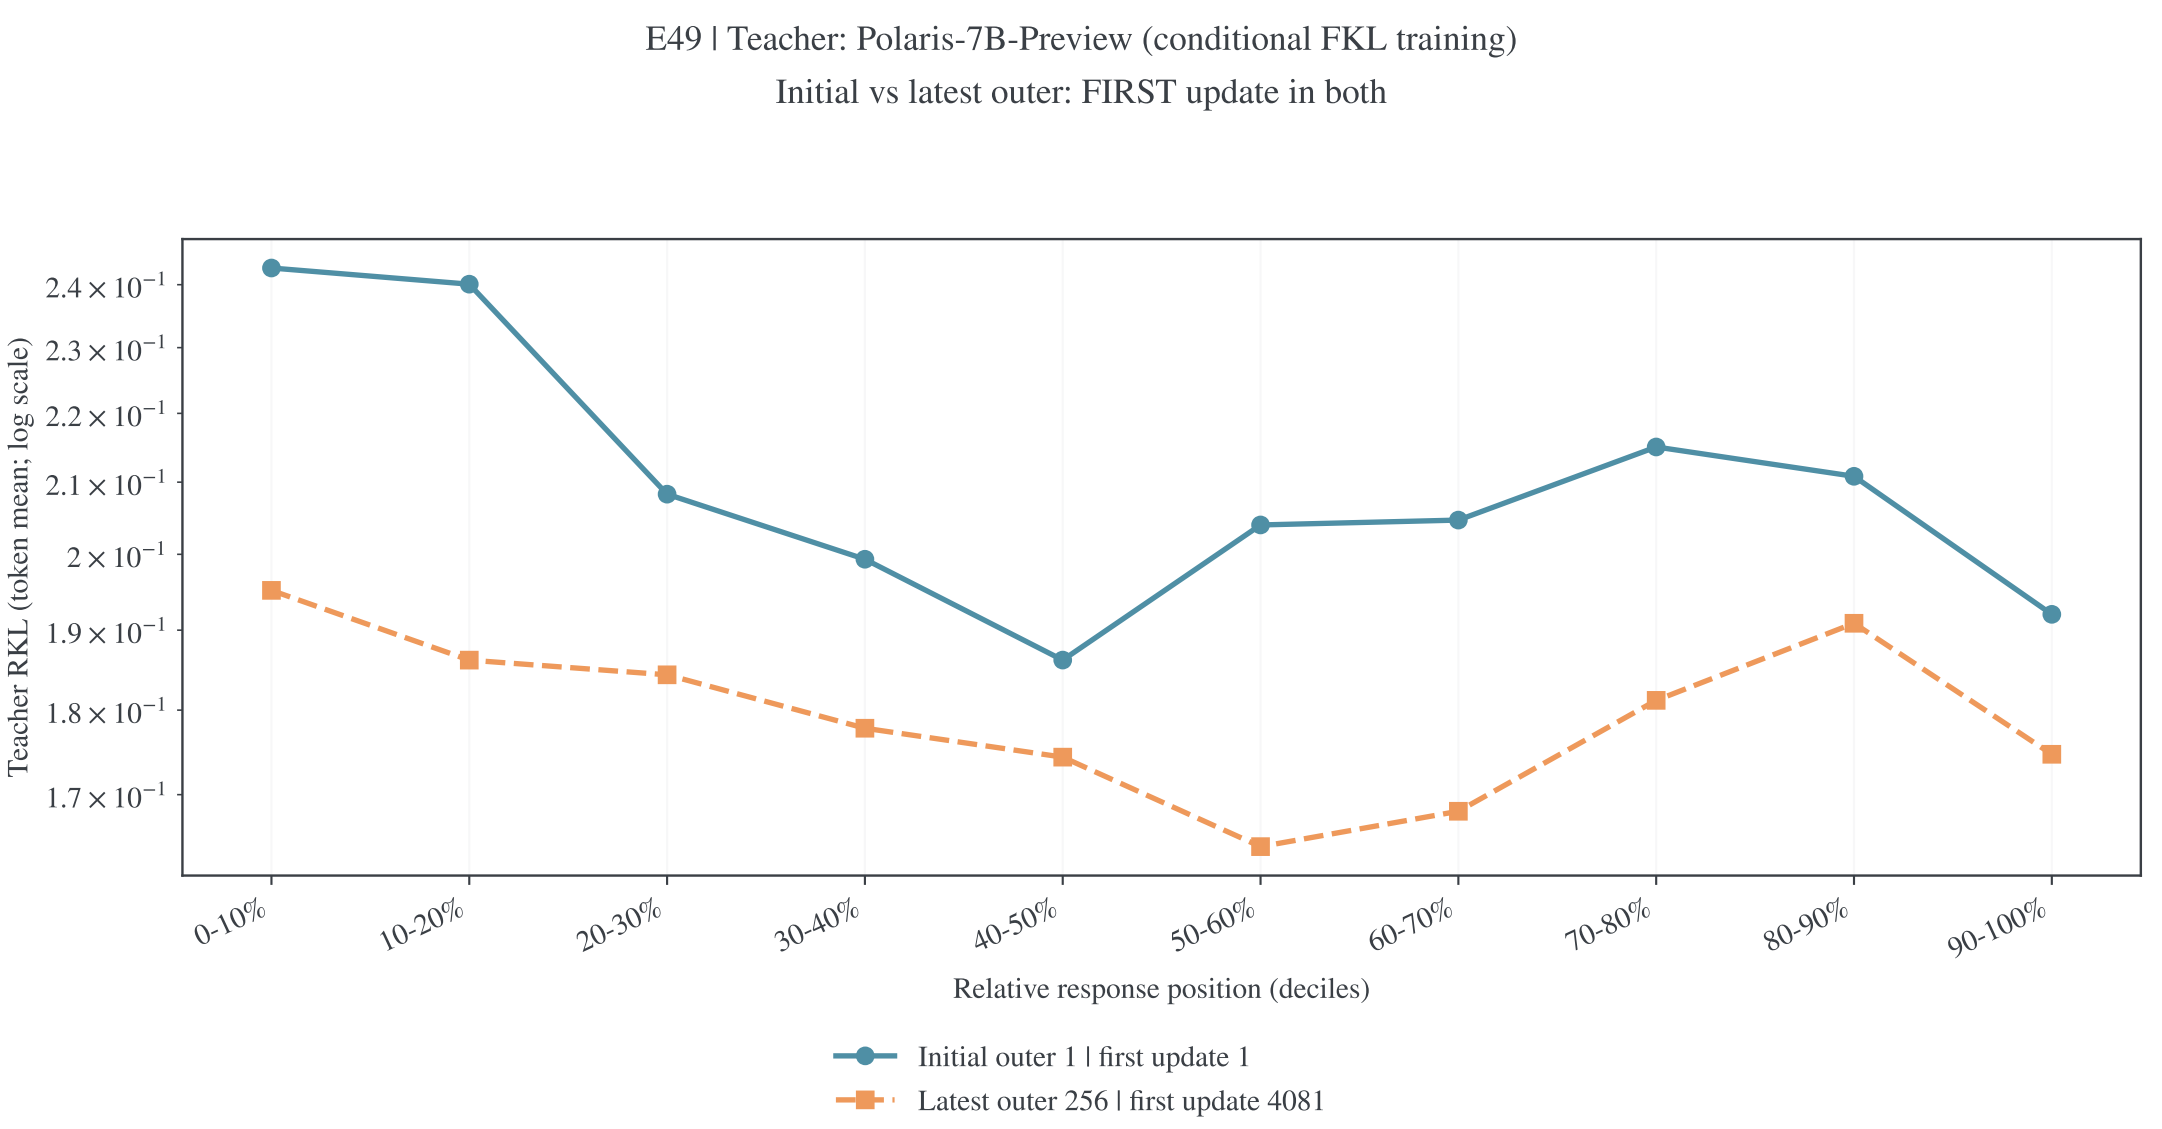
E49 | Teacher: Polaris-7B-Preview (conditional FKL training)
Initial vs latest outer: FIRST update in both
Teacher RKL (token mean; log scale)
2.4×10^{-1}
2.3×10^{-1}
2.2×10^{-1}
2.1×10^{-1}
2×10^{-1}
1.9×10^{-1}
1.8×10^{-1}
1.7×10^{-1}
0-10%
10-20%
20-30%
30-40%
40-50%
50-60%
60-70%
70-80%
80-90%
90-100%
Relative response position (deciles)
Initial outer 1 | first update 1
Latest outer 256 | first update 4081
Both observations precede the first optimizer operation on their own rollout batch; no endpoint averaging.