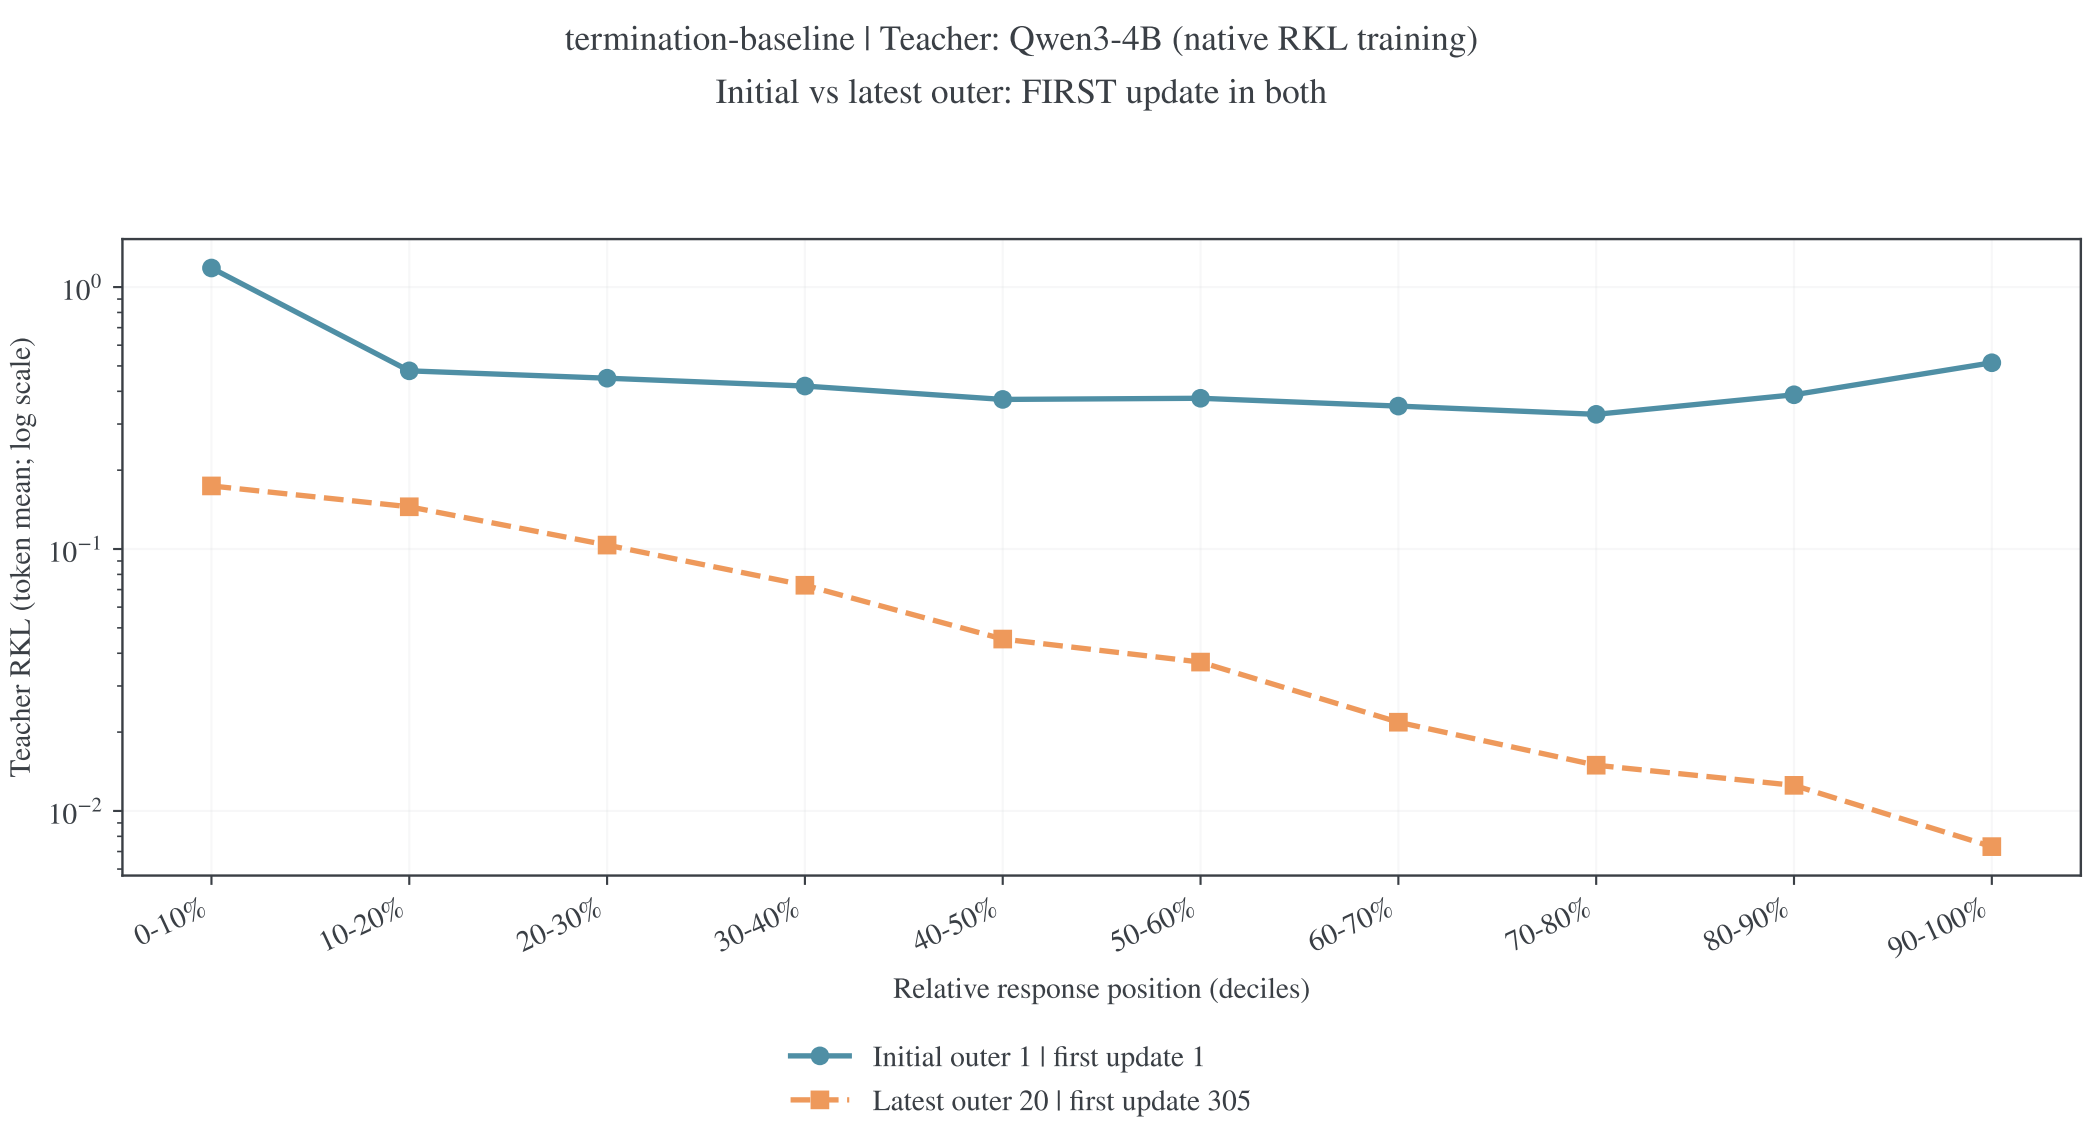
termination-baseline | Teacher: Qwen3-4B (native RKL training)
Initial vs latest outer: FIRST update in both
Teacher RKL (token mean; log scale)
10^{0}
10^{-1}
10^{-2}
0-10%
10-20%
20-30%
30-40%
40-50%
50-60%
60-70%
70-80%
80-90%
90-100%
Relative response position (deciles)
Initial outer 1 | first update 1
Latest outer 20 | first update 305
Both observations precede the first optimizer operation on their own rollout batch; no endpoint averaging.