q3-1.7: same-size checkpoint comparison, 320 updates

RKL / (prompts x 1024)

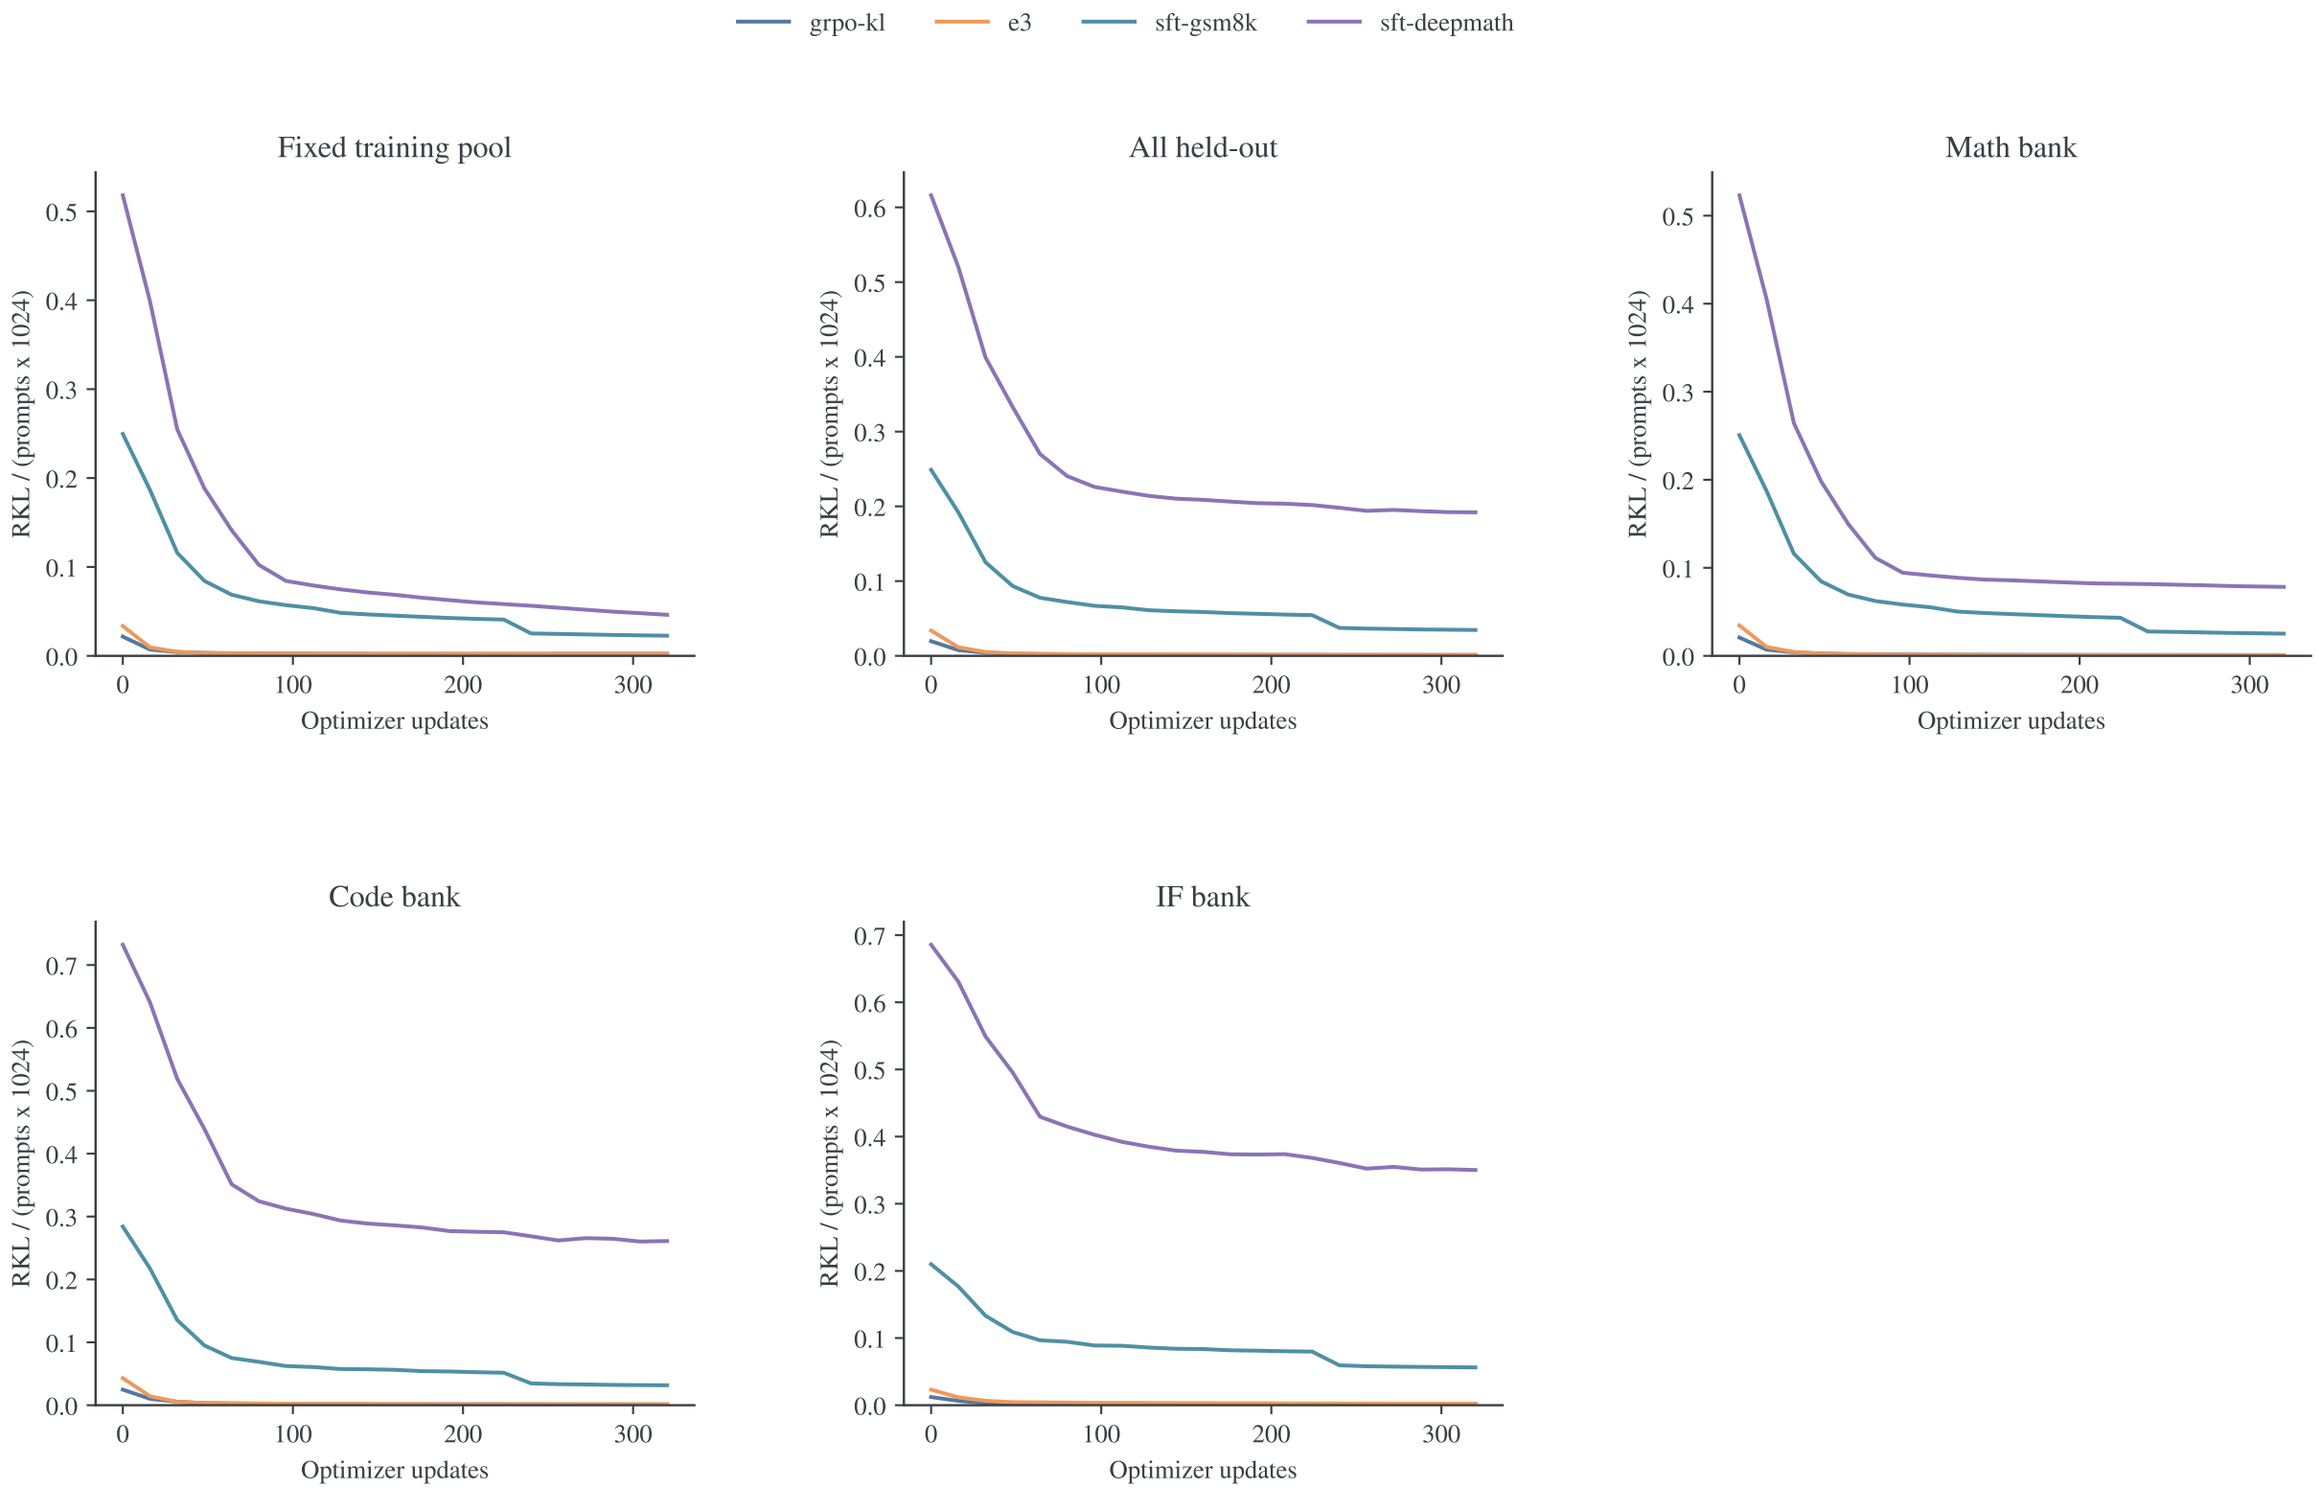

Same student, input mode and 320-update protocol; public checkpoint history and initial KL differ. These are checkpoint comparisons, not controlled RL vs SFT interventions.

First pass only; no repeated-batch round. Same size does not establish direct ancestry.

post-trained dual-mode; thinking input; native_old_thinking