Qwen3-1.7B-GRPO-KL-Math -> Qwen3-1.7B

First-pass training loss and within-batch progress

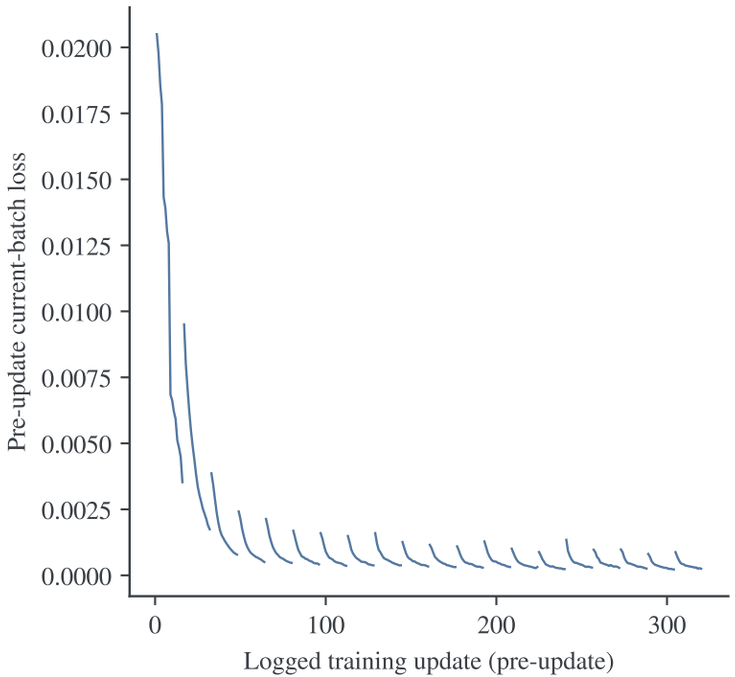

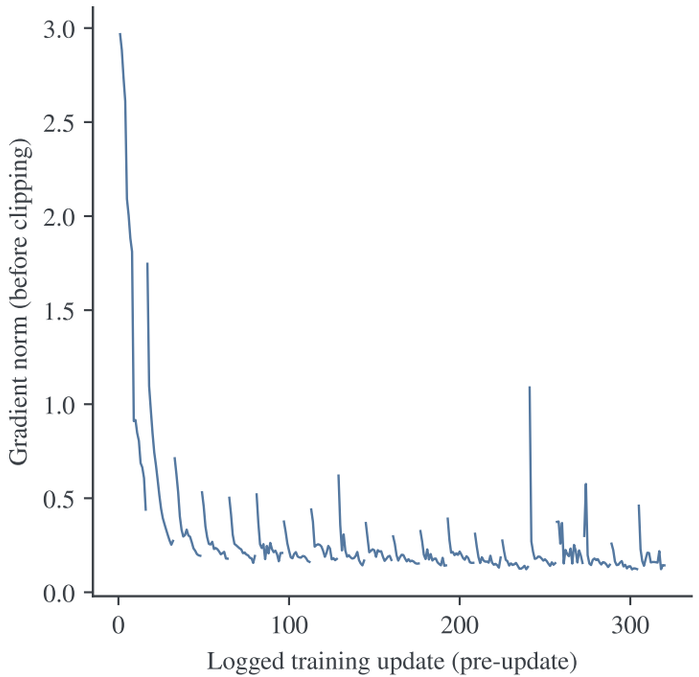

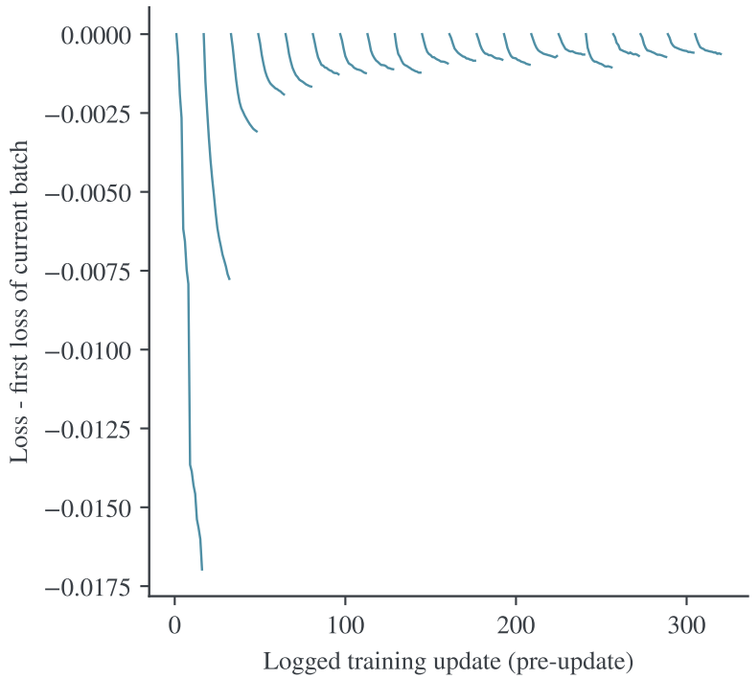

Batch16 x reuse16; segments stop at switches. A logged loss at step t scores theta_(t-1), not the fixed theta_t snapshot.

First pass only; no repeated-batch round. Same size does not establish direct ancestry.

post-trained dual-mode; thinking input; native_old_thinking