R1-1.5B | RKL along training (log-log)

(a) Math

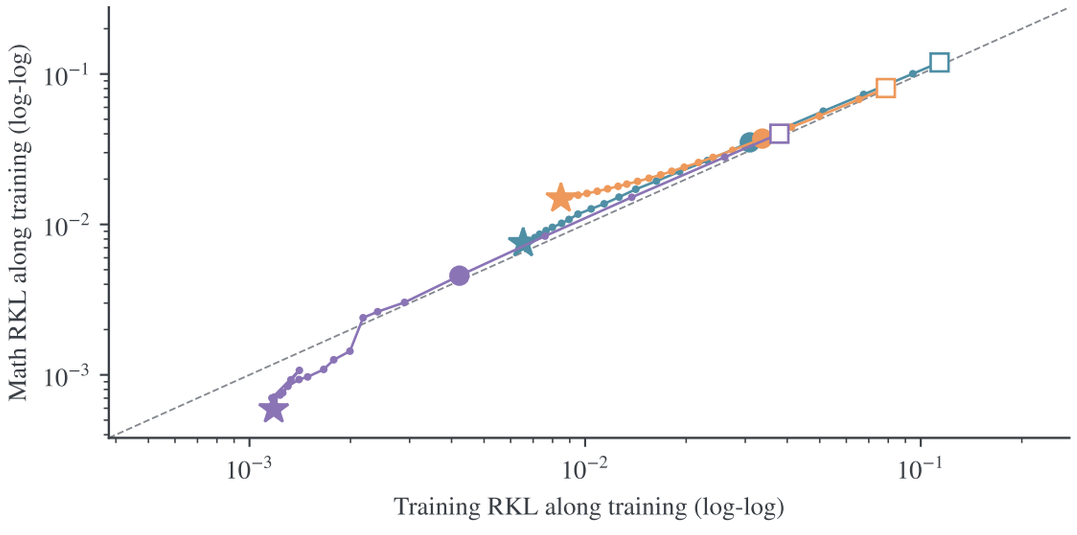

(b) Code

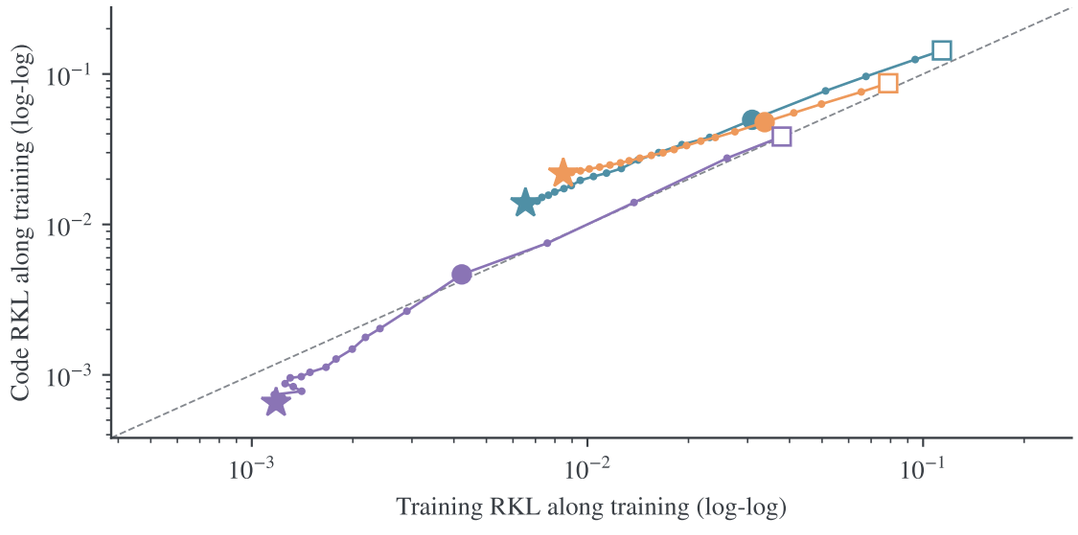

(c) IF

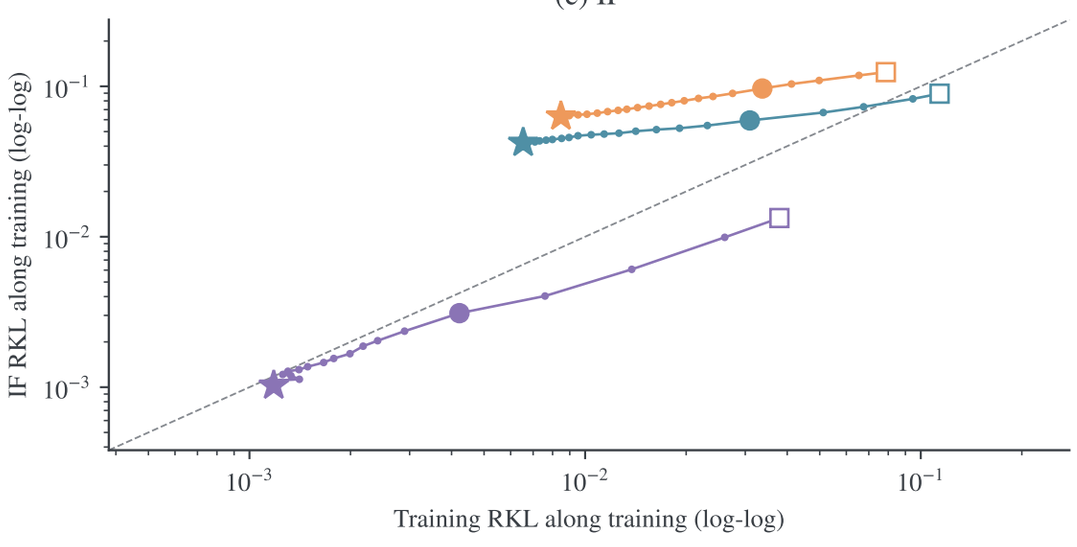

(d) Avg

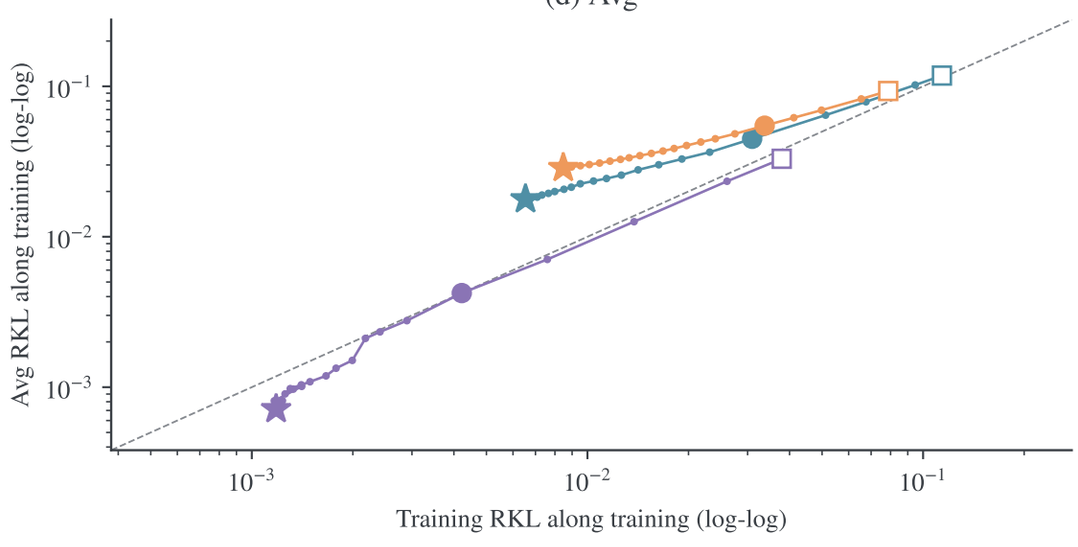

JustRL (math RL) MiniMath (math SFT) DeepScaleR (math RL) Equality

Each dot = one snapshot, every 16 updates; square: 0; large circle: 64; star: 320.

Lines follow time within one run; snapshots are dependent, so no cross-run correlation is estimated.