Llama-3.2-3B-Instruct | Training pool vs. held-out

Absolute RKL decrease (nats / H1024)

(a) Training vs. Math

n=10; r=0.98; rho=0.90

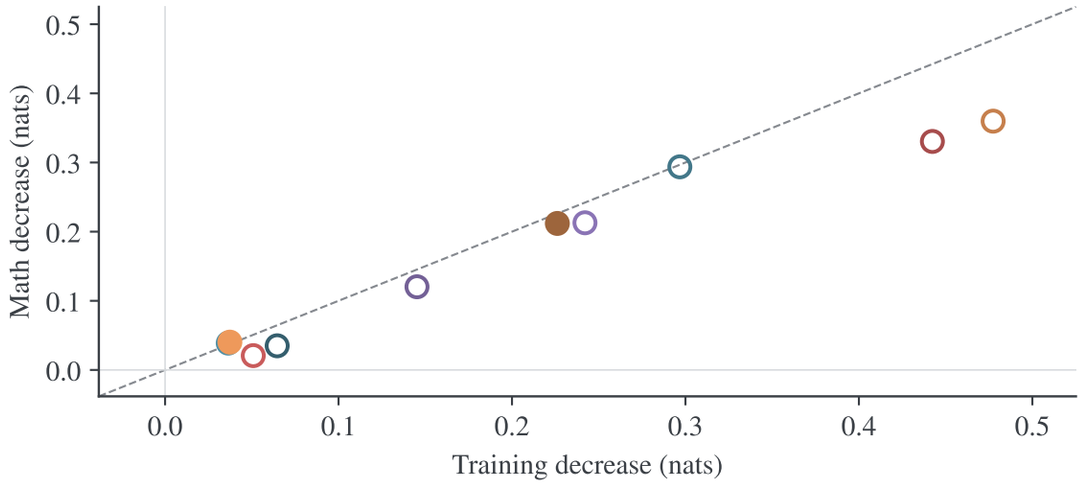

(b) Training vs. Code

n=10; r=0.84; rho=0.88

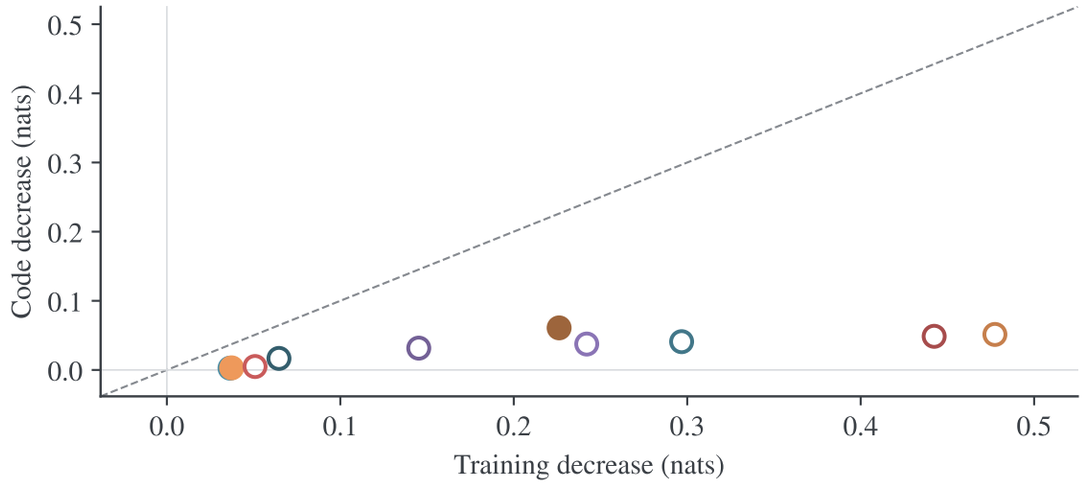

(c) Training vs. IF

n=10; r=0.70; rho=0.82

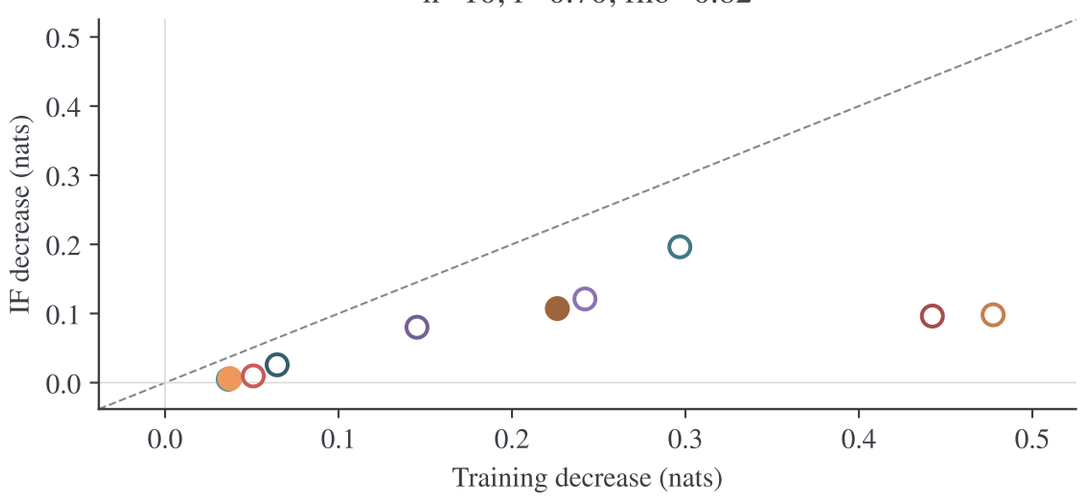

(d) Training vs. Avg

n=10; r=0.95; rho=0.94

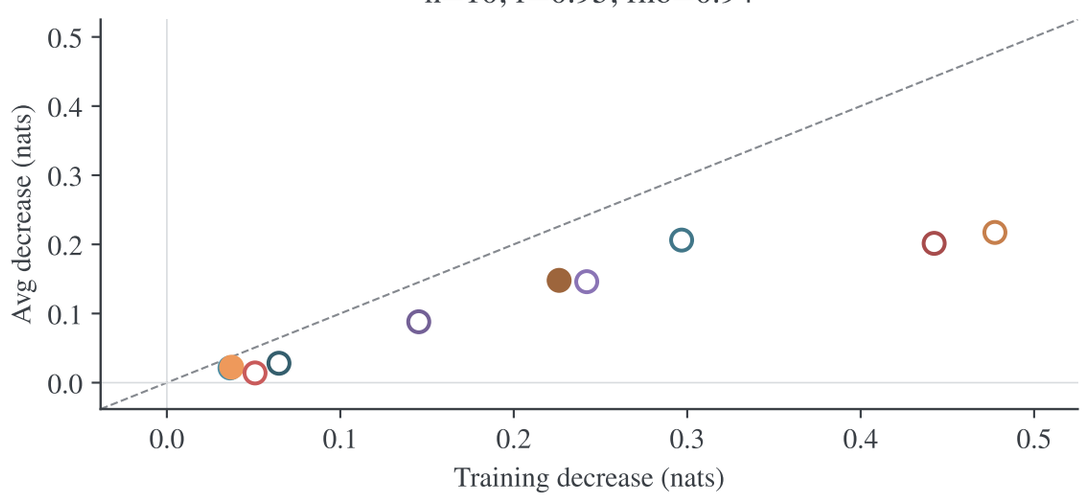

Co-rewarding-I-Llama-3.2-3B-Instruct-MATH [320]

GT-Llama-3.2-3B-Instruct-MATH [320]

Llama-3.1-8B-Instruct / Unsloth [320]

Llama-3.1-8B-SimpleRL-Zoo [320]

Llama-3.1-Nemotron-Nano-8B-v1 [320]

Llama-3.1-Tulu-3-8B [320]

Llama-3.1-Tulu-3-8B-DPO [320]

Llama-3.1-Tulu-3-8B-SFT [320]

Llama-3.2-1B-Instruct / Unsloth [320]

Llama-3.2-3B-Instruct-arithmetic-full [320]

Equal decrease

One point = one run endpoint; triangle: 64 updates; circle: 320; filled: same size; open: cross size.

Fixed-H RKL; dashed: equality; correlations pool all displayed endpoints (mixed budgets/sizes).