Qwen3-1.7B / native Thinking-2507 | Training pool vs. held-out

Final RKL (log-log axes; correlations on log10 RKL)

(a) Training vs. Math

n=1; correlation unavailable

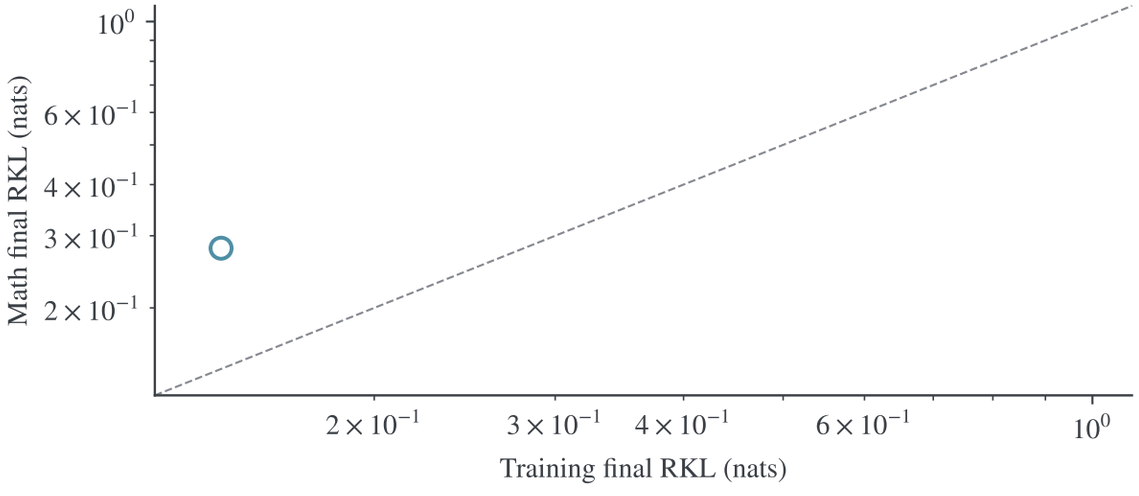

(b) Training vs. Code

n=1; correlation unavailable

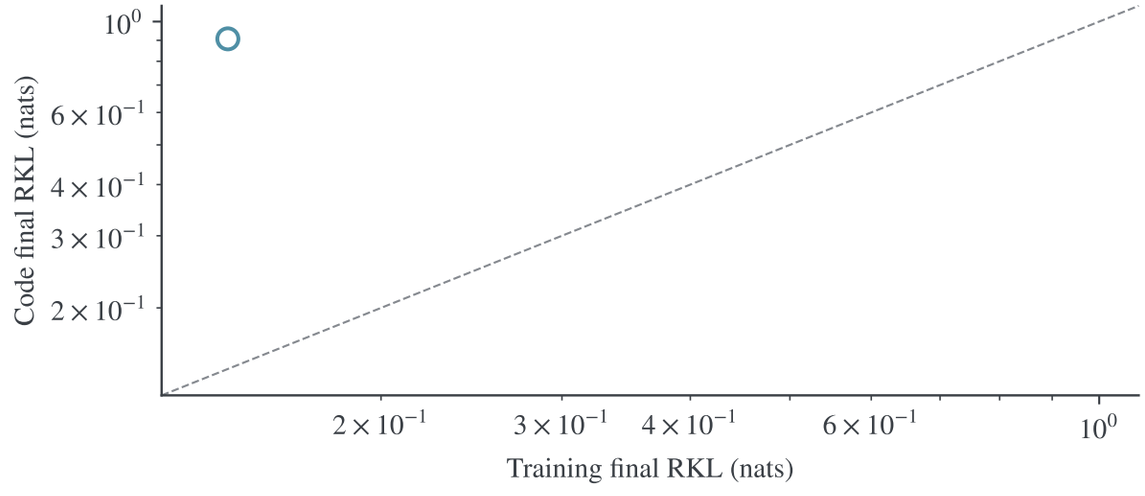

(c) Training vs. IF

n=1; correlation unavailable

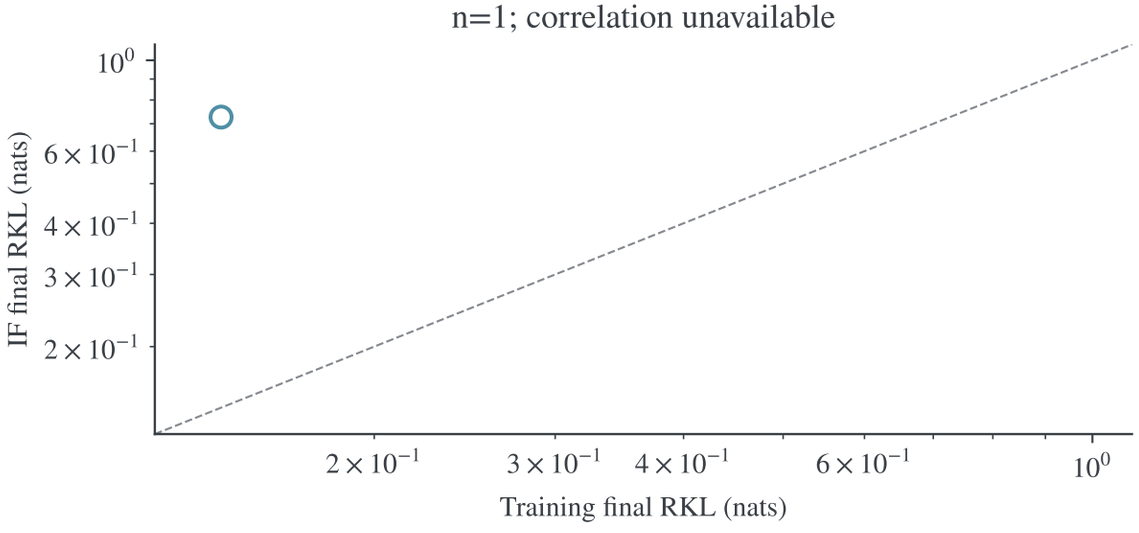

(d) Training vs. Avg

n=1; correlation unavailable

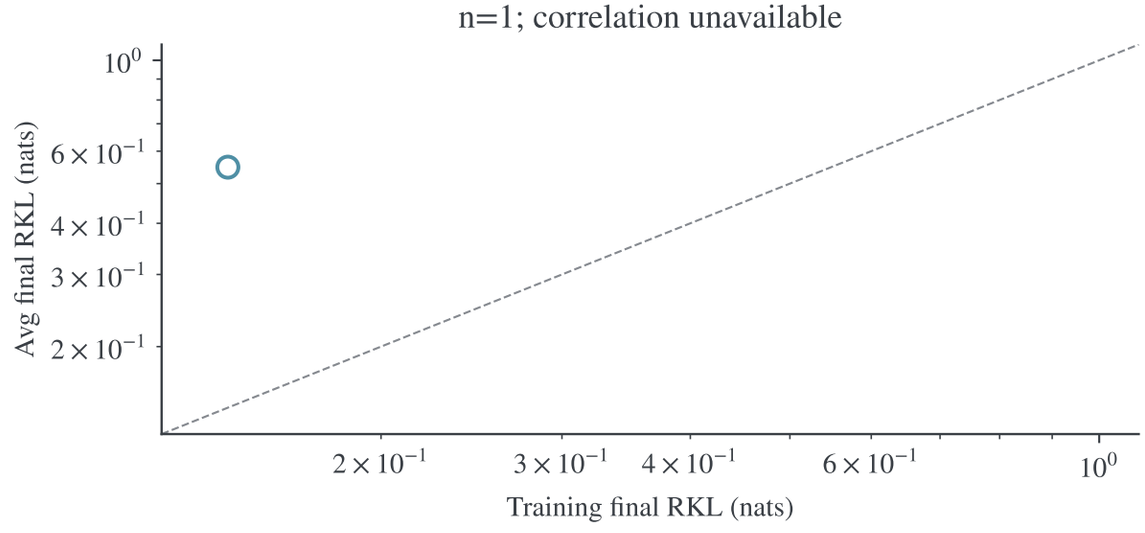

Qwen3-4B-Thinking-2507 [320] Equal final RKL

One point = one run endpoint; triangle: 64 updates; circle: 320; filled: same size; open: cross size.

Fixed-H RKL; dashed: equality; correlations pool all displayed endpoints (mixed budgets/sizes).