Qwen3-1.7B / non-thinking | Held-out domain vs. domain

Relative RKL decrease (%)

(a) Math vs. Code

n=10; r=0.87; rho=0.79

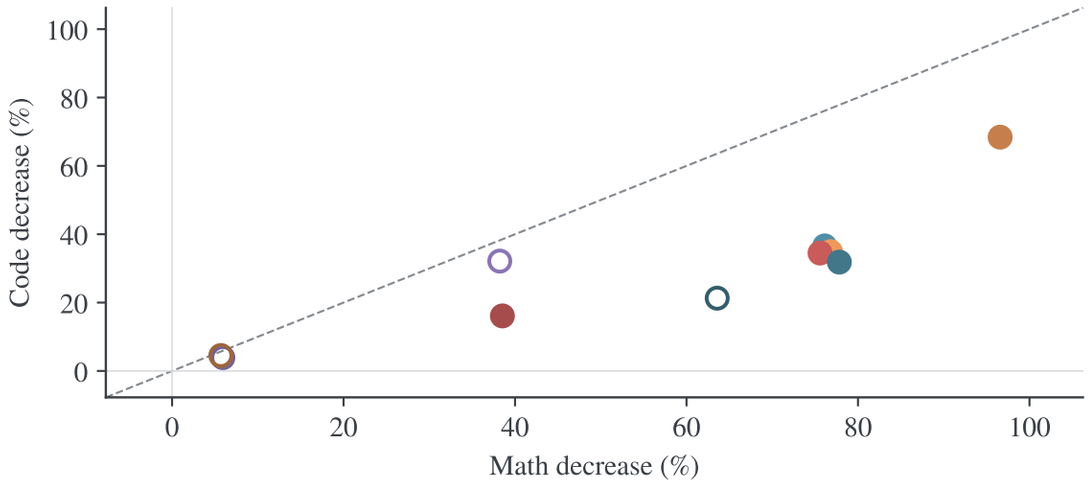

(b) Math vs. IF

n=10; r=0.84; rho=0.93

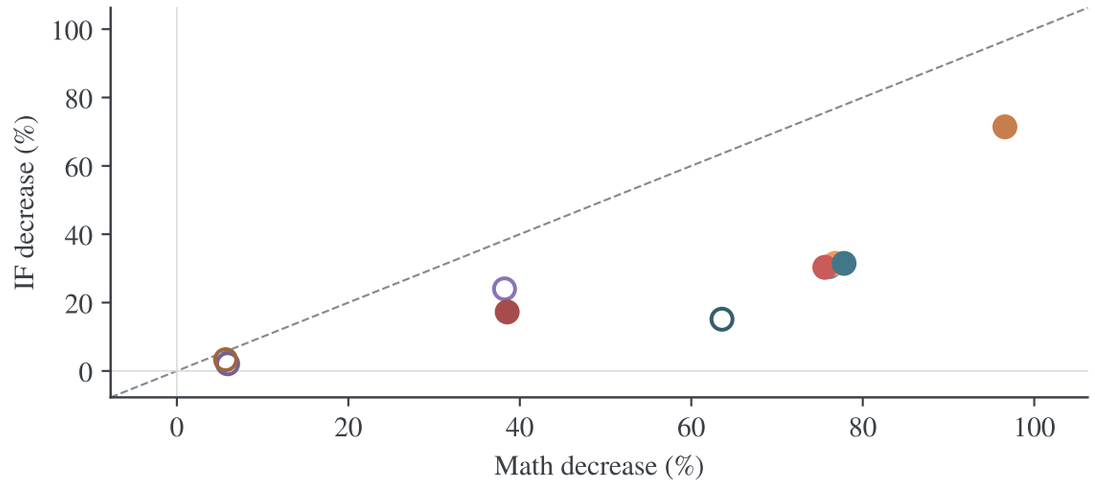

(c) Code vs. IF

n=10; r=0.98; rho=0.89

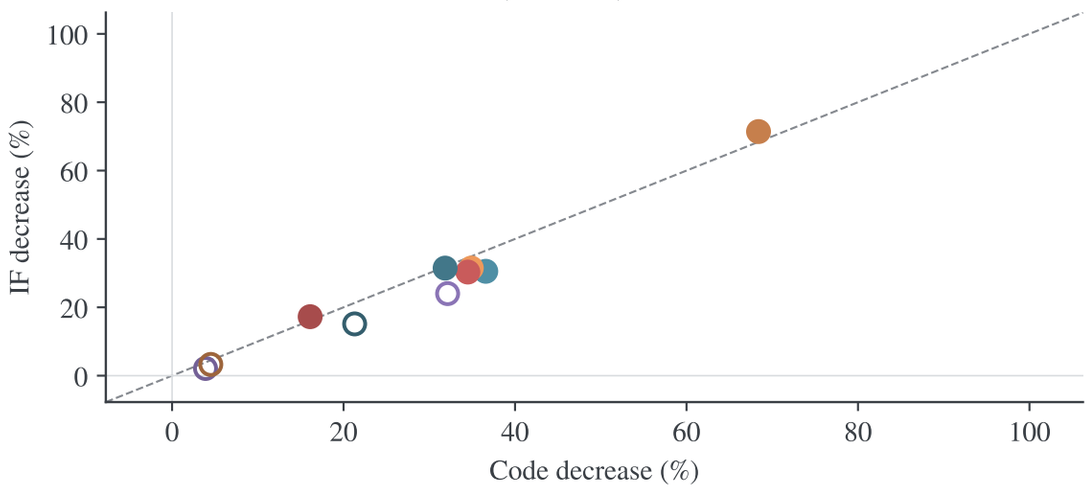

- Co-rewarding-I-Qwen3-1.7B-Base-MATH [320]
- Co-rewarding-II-Qwen3-1.7B-Base-MATH [320]
- GT-Qwen3-1.7B-Base-MATH [320]
- Qwen3-0.6B [320]
- Qwen3-1.7B-Base-OPD / Thinking-Space [320]
- Qwen3-1.7B-Math / Samsung SFT [320]
- Qwen3-1.7B-SFT / Thinking-Space [320]
- Qwen3-4B [320]
- Qwen3-4B-Base-GRPO / Thinking-Space [320]
- Qwen3-8B [320]
- Equal decrease

One point = one run endpoint; triangle: 64 updates; circle: 320; filled: same size; open: cross size.

Fixed-H RKL; dashed: equality; correlations pool all displayed endpoints (mixed budgets/sizes).