R1-1.5B | Training pool vs. held-out

Final RKL (log-log axes; correlations on log10 RKL)

(a) Training vs. Math

n=9; r=0.99; rho=0.98

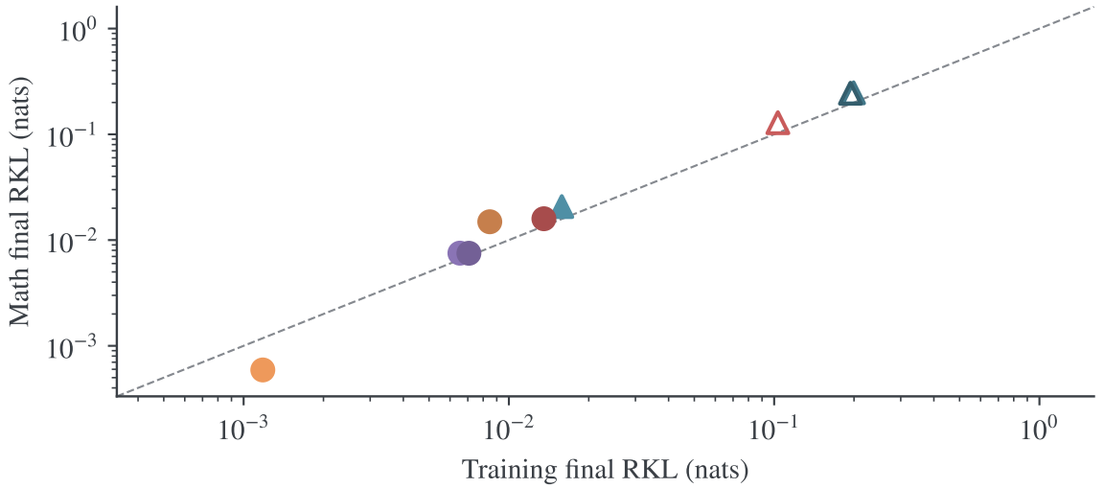

(b) Training vs. Code

n=9; r=0.90; rho=0.97

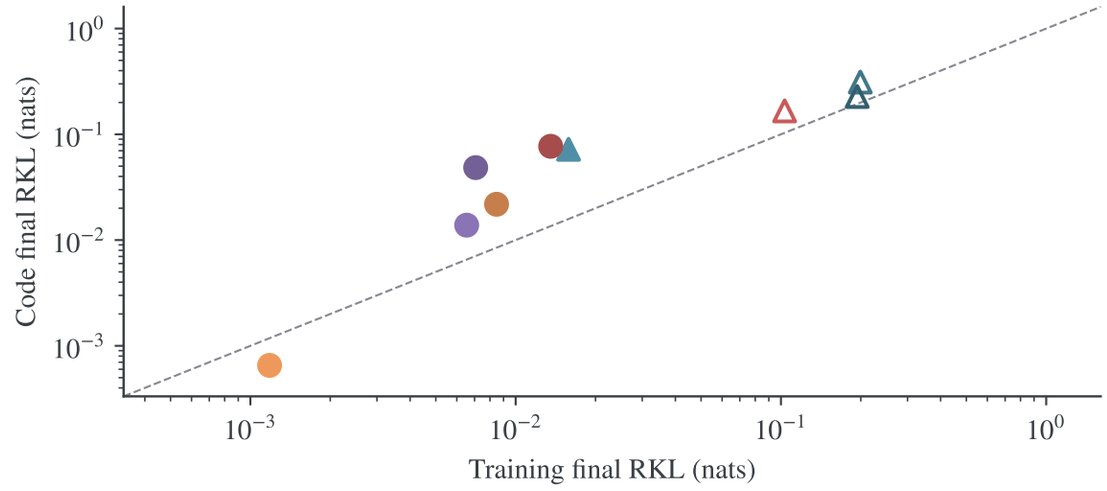

(c) Training vs. IF

n=9; r=0.87; rho=0.82

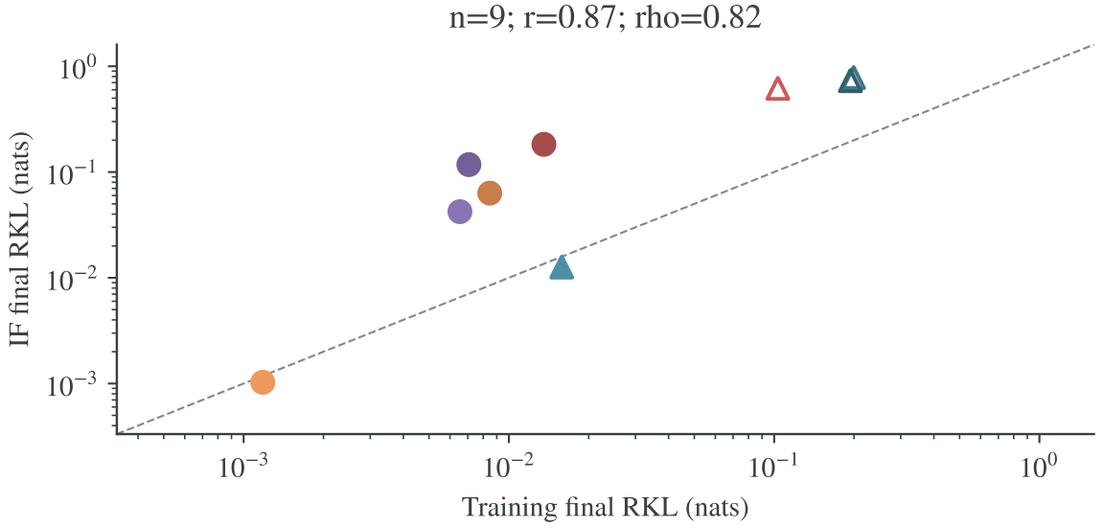

(d) Training vs. Avg

n=9; r=0.94; rho=0.92

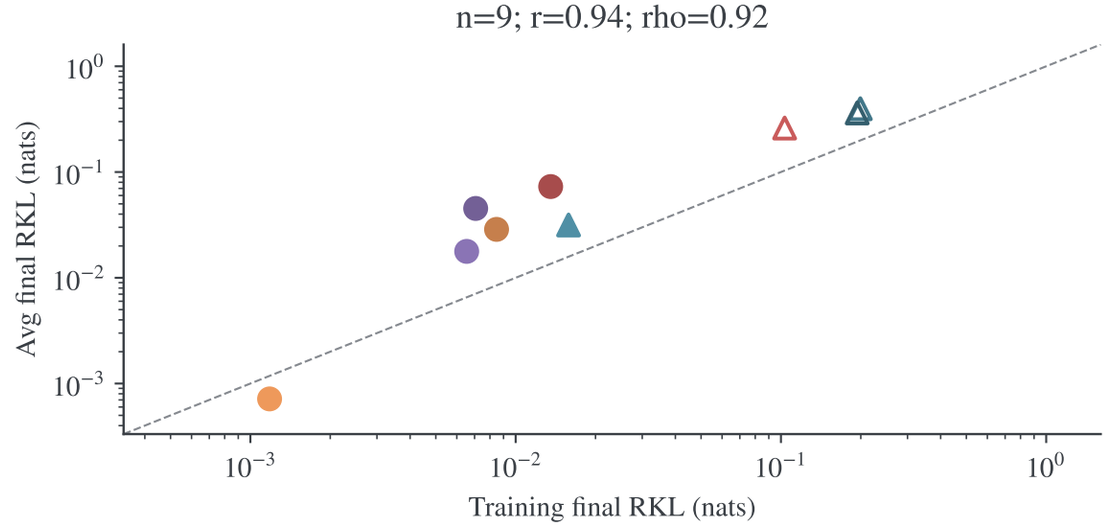

DeepCoder-1.5B-Preview [64]
DeepScaleR-1.5B-Preview [320]
DeepSeek-R1-Distill-Qwen-7B [64]
JustRL-DeepSeek-1.5B [320]
Light-R1-7B-DS [64]
MiniMath-R1-1.5B [320]
Nemotron-Research-Reasoning-Qwen-1.5B / BroRL [320]
Nemotron-Research-Reasoning-Qwen-1.5B v1 [320]
Skywork-OR1-Math-7B [64]
Equal final RKL

One point = one run endpoint; triangle: 64 updates; circle: 320; filled: same size; open: cross size.

Fixed-H RKL; dashed: equality; correlations pool all displayed endpoints (mixed budgets/sizes).