R1-1.5B | RKL along training (log-log)

Same size: solid = RL, dashed = added SFT, dotted = other controls

(a) Math

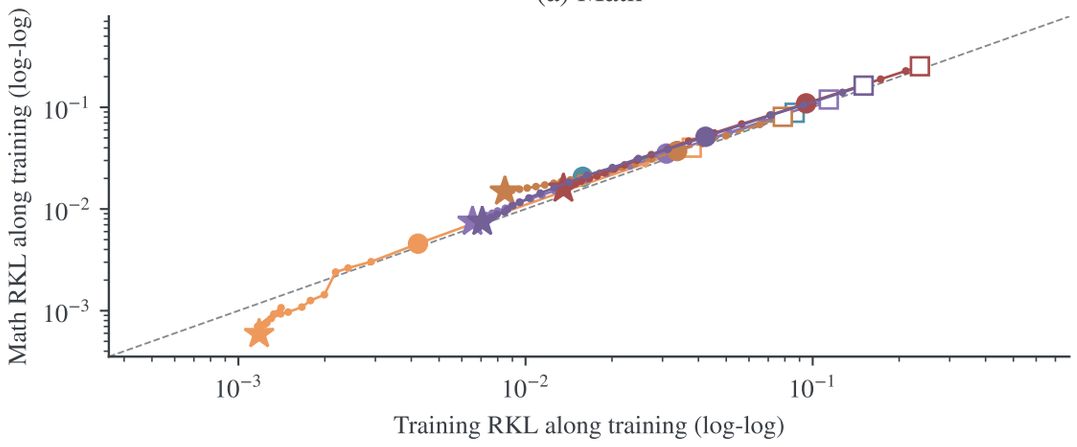

(b) Code

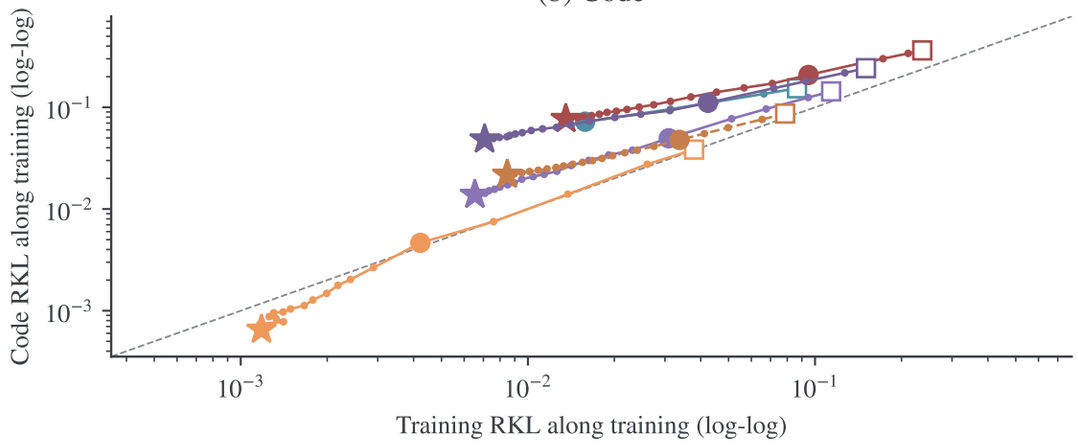

(c) IF

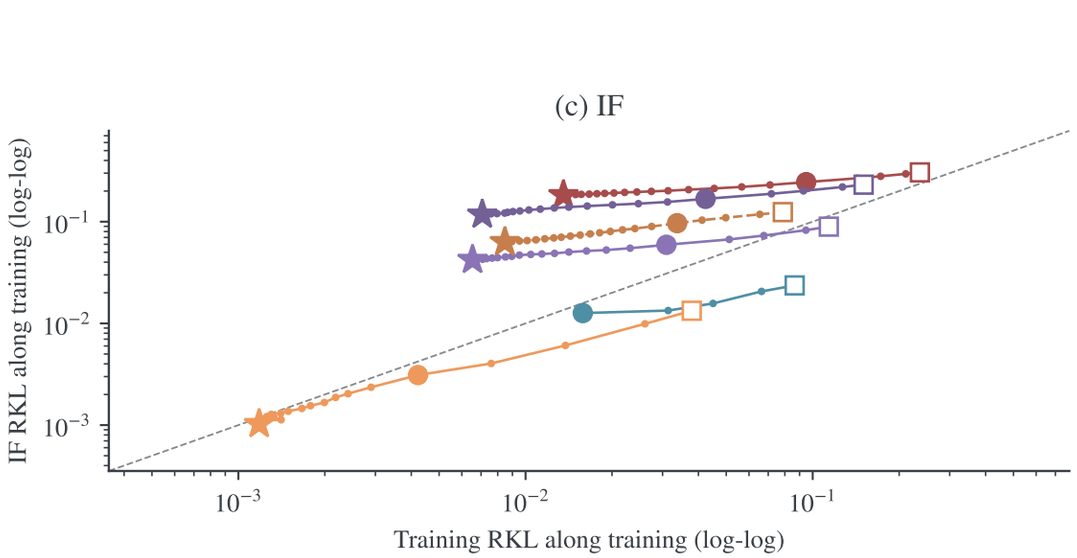

(d) Avg

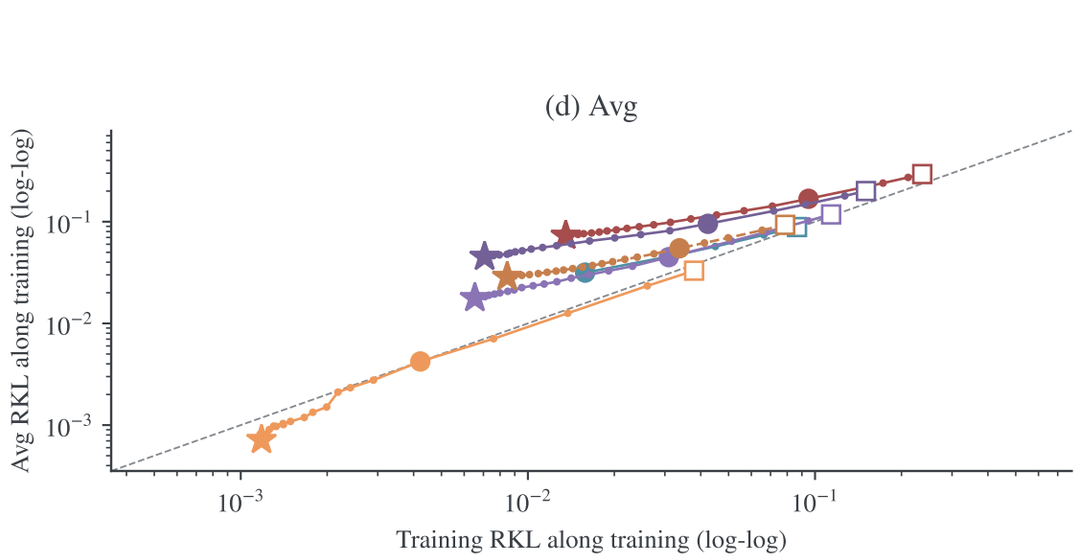

deepcoder [64; RL]
deepscaler [320; RL]
justrl [320; RL]
minimath [320; SFT]
brorl [320; RL]
nemotron [320; RL]
Equality

Each dot = one snapshot, every 16 updates; square: 0; large circle: 64; star: 320.

Lines follow time within one run; snapshots are dependent, so no cross-run correlation is estimated.