E39: Student Qwen3-1.7B-Base ← Teacher Qwen3-8B

Thinking OFF (enable_thinking=False); Full-RKL, reuse16, batch64, LR 1e-6

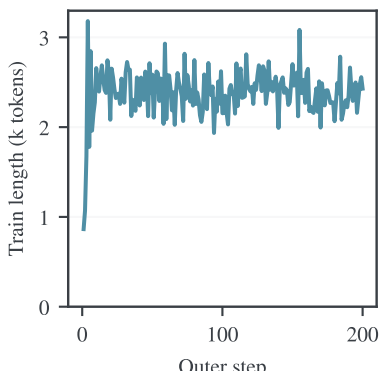

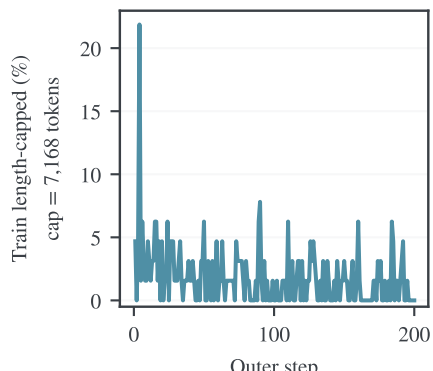

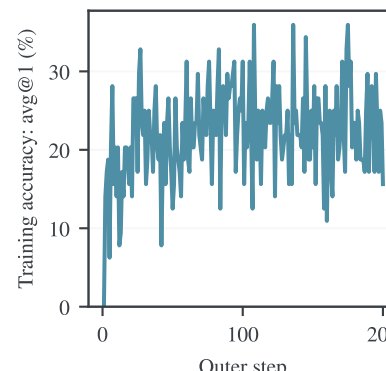

Training cap=7,168 tokens; one response/question (avg@1).

DAPO batches change each step; correctness is diagnostic, outside the loss.