Student: unsloth/Llama-3.2-3B-Instruct | non-thinking / shared off input

Equal budget: 0-320 updates

Fixed training pool A-T (320)

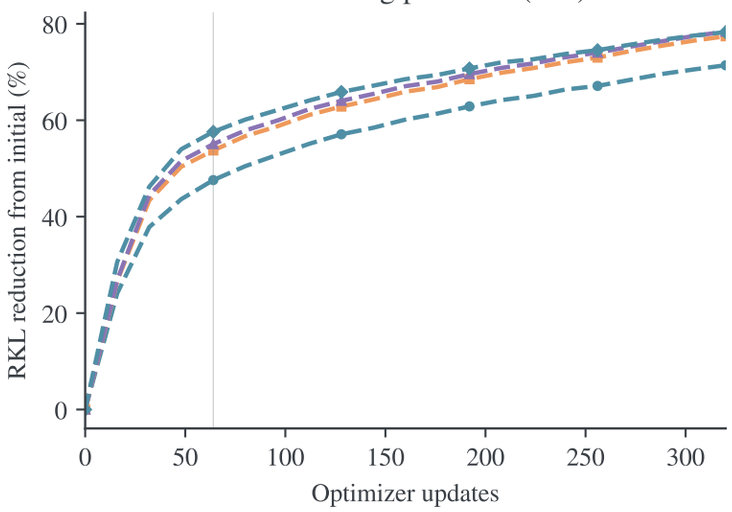

Old training batches A-D (64)

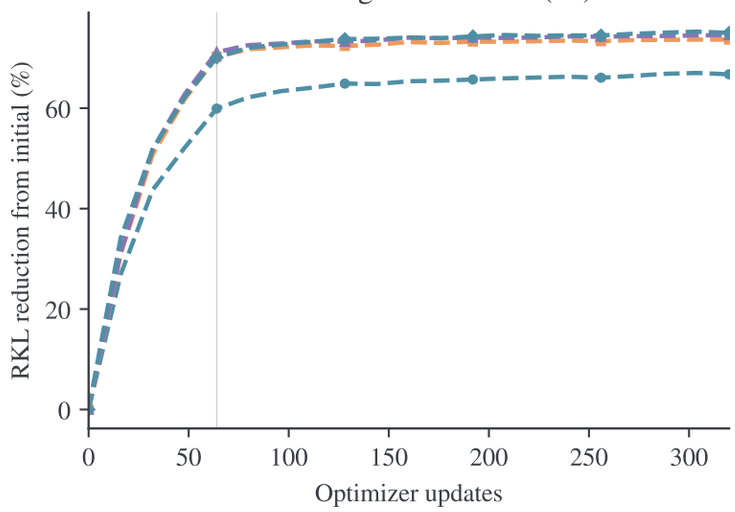

All held-out (128)

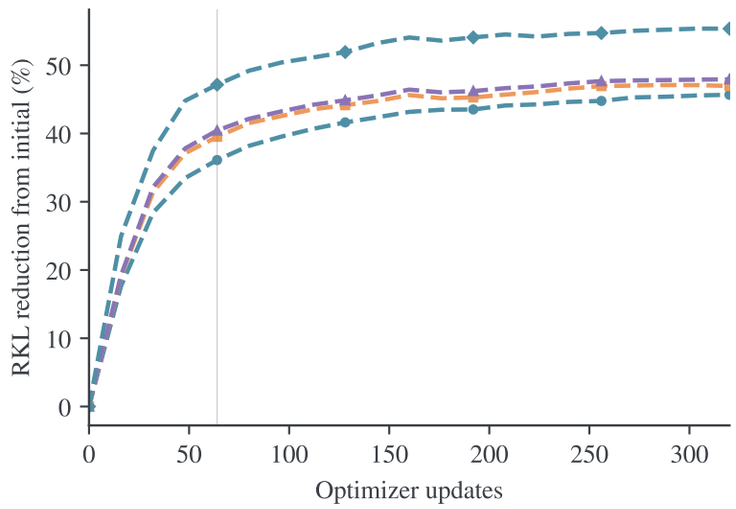

Held-out Math (64)

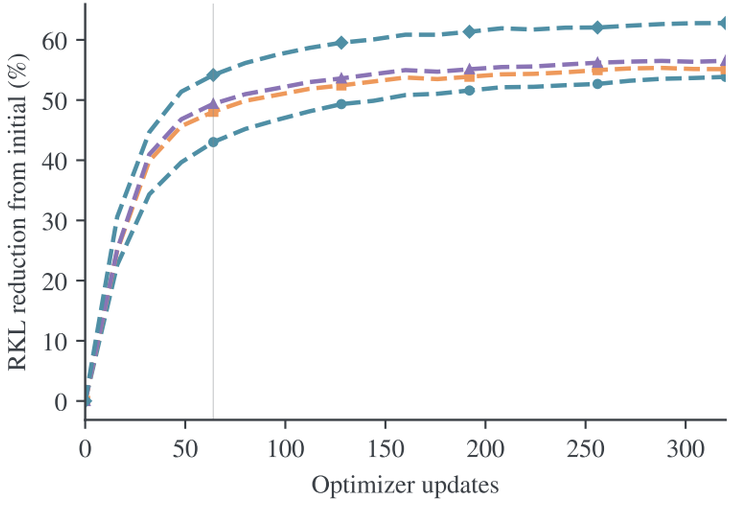

Held-out Code (32)

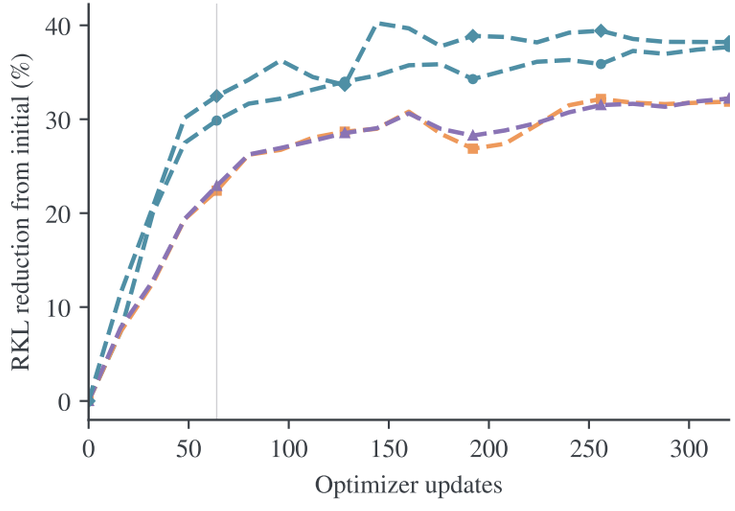

Held-out IF (32)

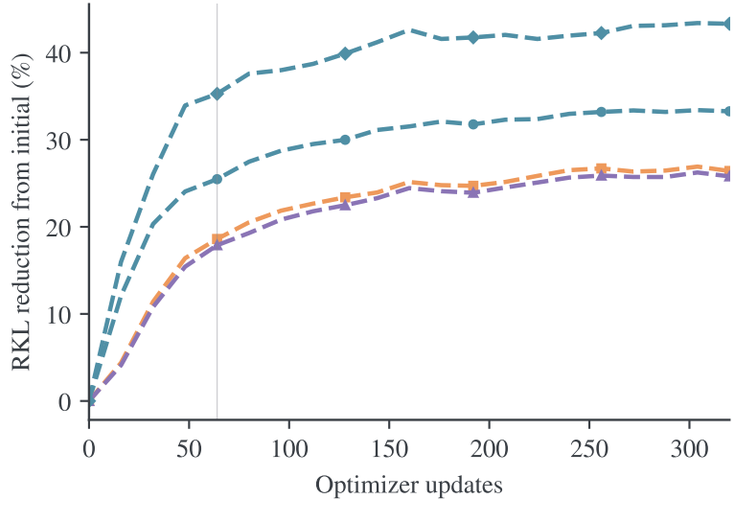

Llama-3.1-Tulu-3-8B-SFT
Llama-3.1-Tulu-3-8B-DPO
Llama-3.1-Tulu-3-8B / final RLVR
Llama-3.1-8B-SimpleRL-Zoo / input-mismatch control

Fixed S0; H1024; one training seed. Full training pool includes not-yet-visited batches early in training.

Solid: same size; dashed: different size. Old A-D is a subset of the training pool. Held-out pooling weights Math/Code/IF by 64/32/32 prompts.