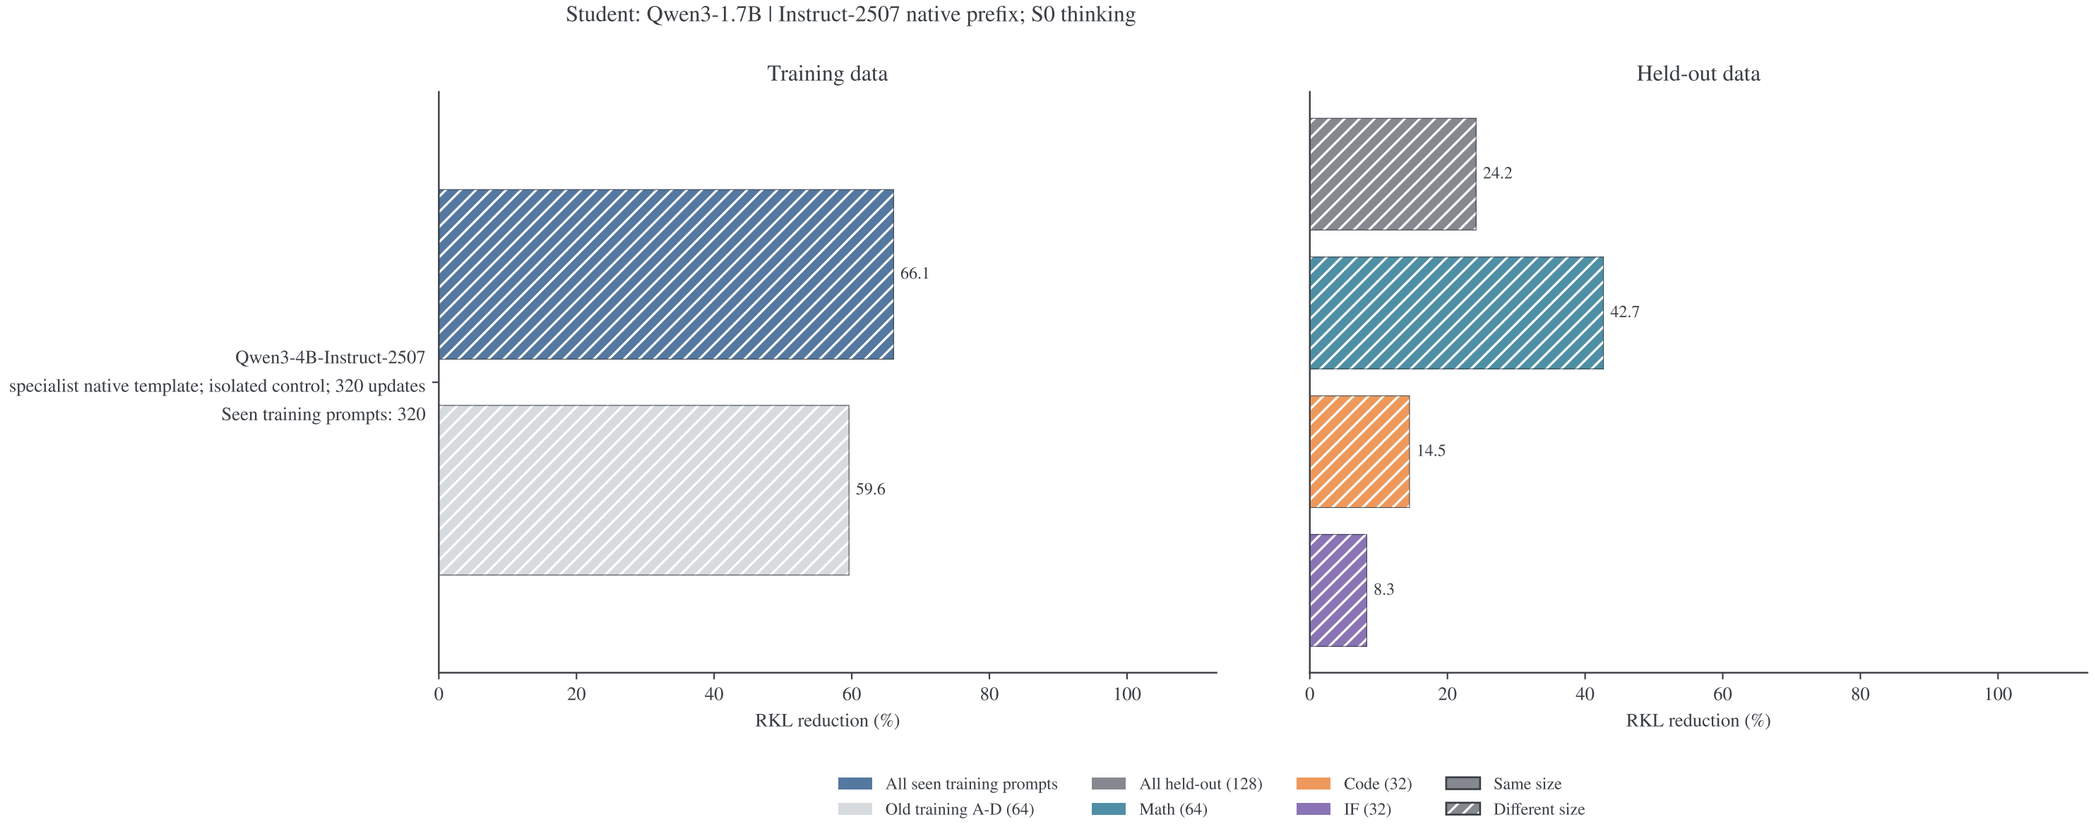
Student: Qwen3-1.7B | Instruct-2507 native prefix; S0 thinking
Training data
Held-out data
Qwen3-4B-Instruct-2507
specialist native template; isolated control; 320 updates
Seen training prompts: 320
66.1
59.6
24.2
42.7
14.5
8.3
0
20
40
60
80
100
RKL reduction (%)
All seen training prompts
Old training A-D (64)
All held-out (128)
Math (64)
Code (32)
IF (32)
Same size
Different size
Rows sorted by pooled held-out RKL reduction (Math64 + Code32 + IF32). Shared scale across panels.
Training = prompts actually used for updates; A-D is a subset, not an extra dataset. Frozen S0; H1024.