Student: Qwen3-1.7B | post-trained dual-mode; non-thinking input

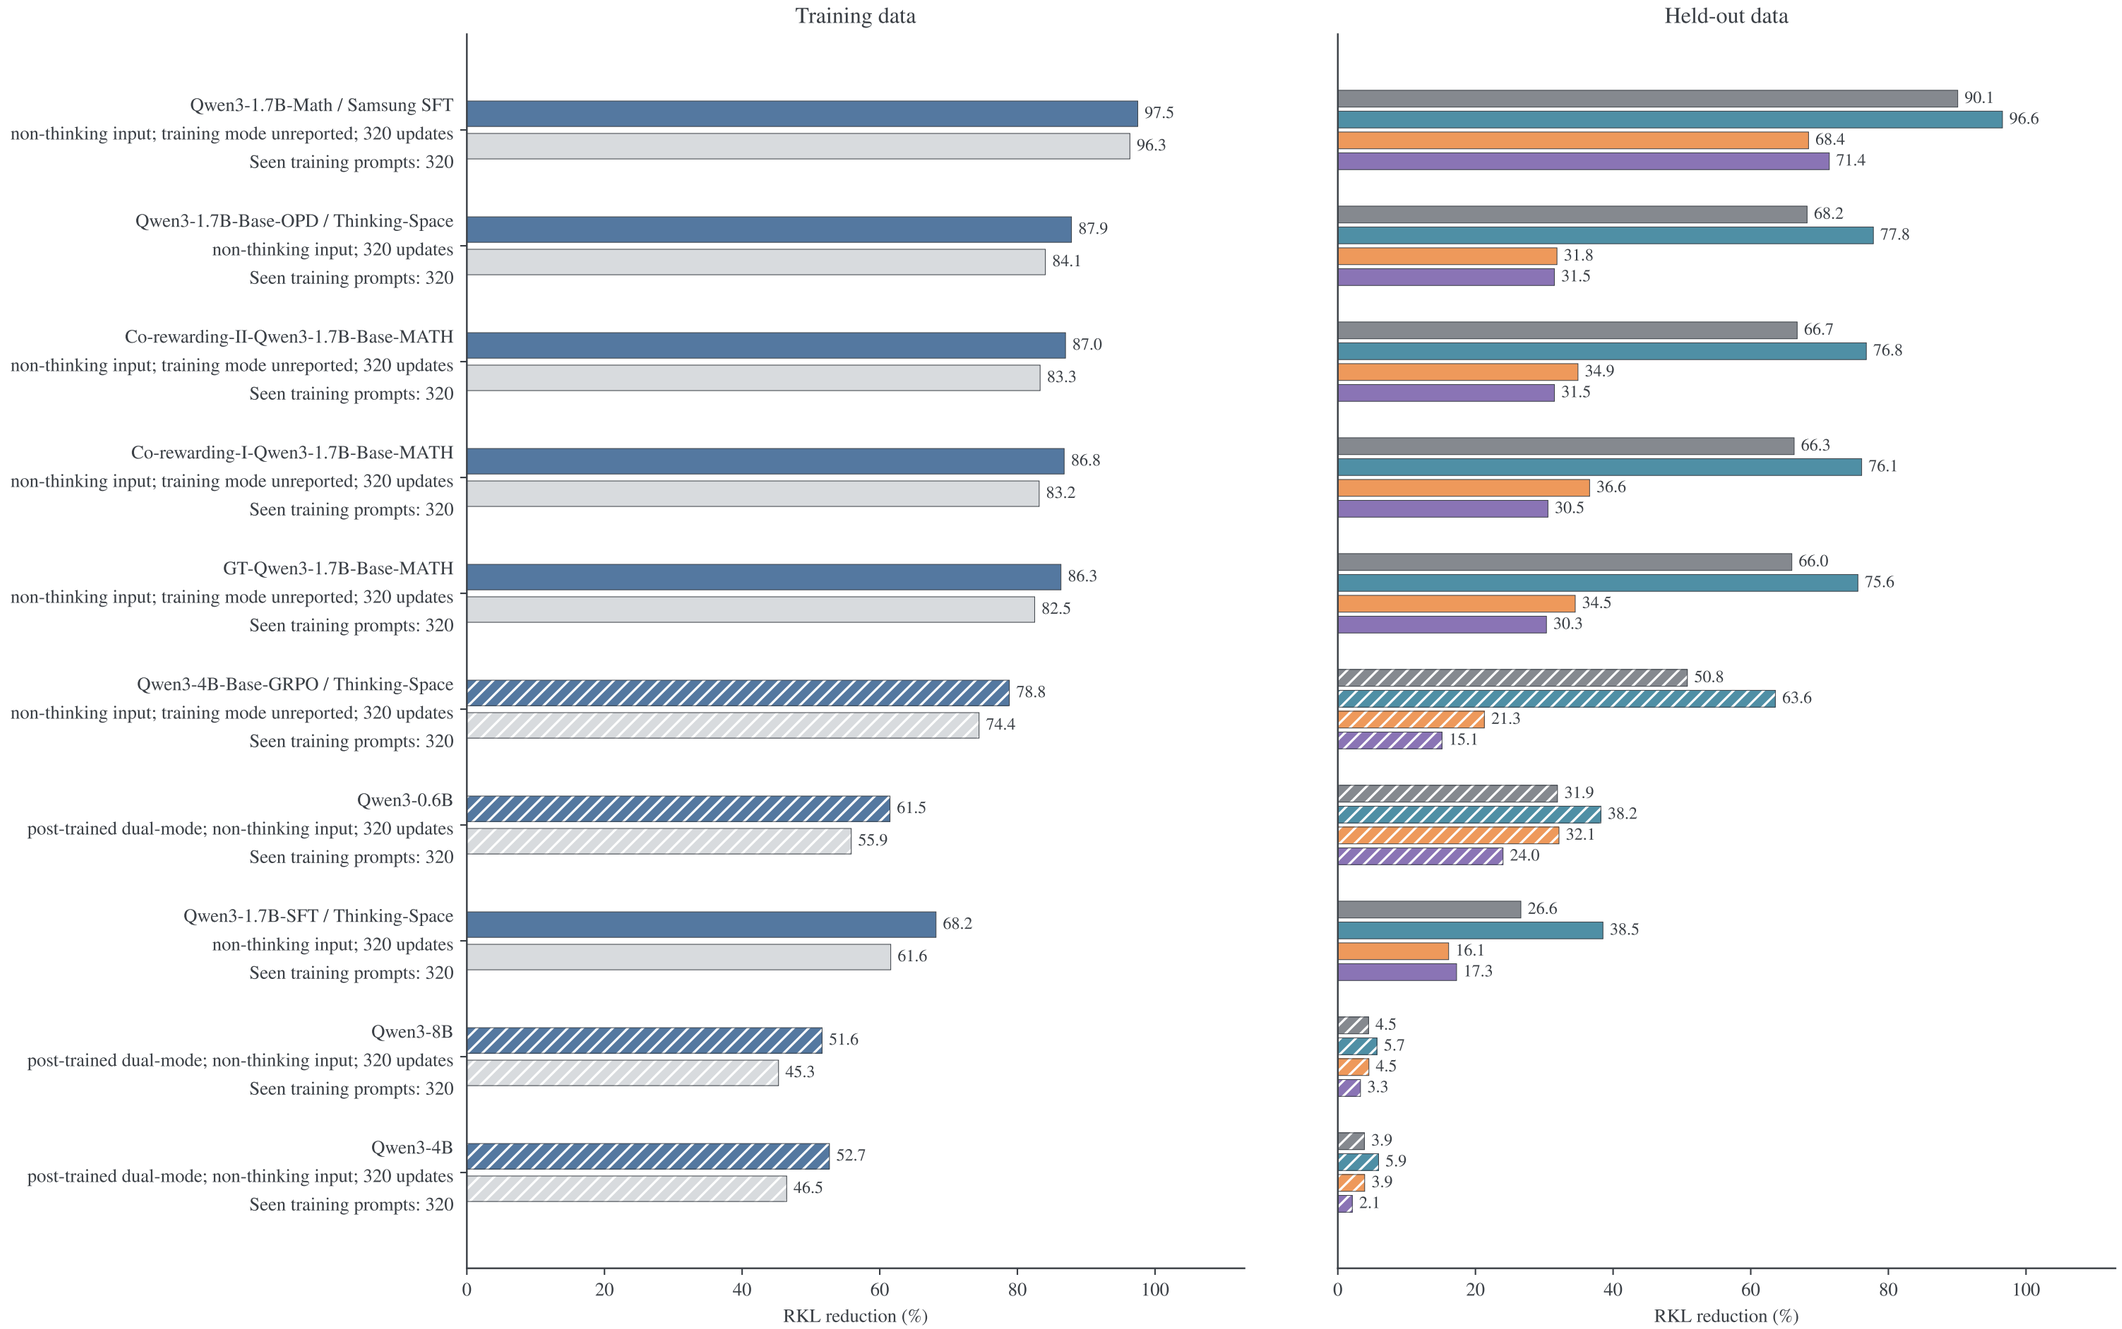

Rows sorted by pooled held-out RKL reduction (Math64 + Code32 + IF32). Shared scale across panels.
Training = prompts actually used for updates; A-D is a subset, not an extra dataset. Frozen S0; H1024.