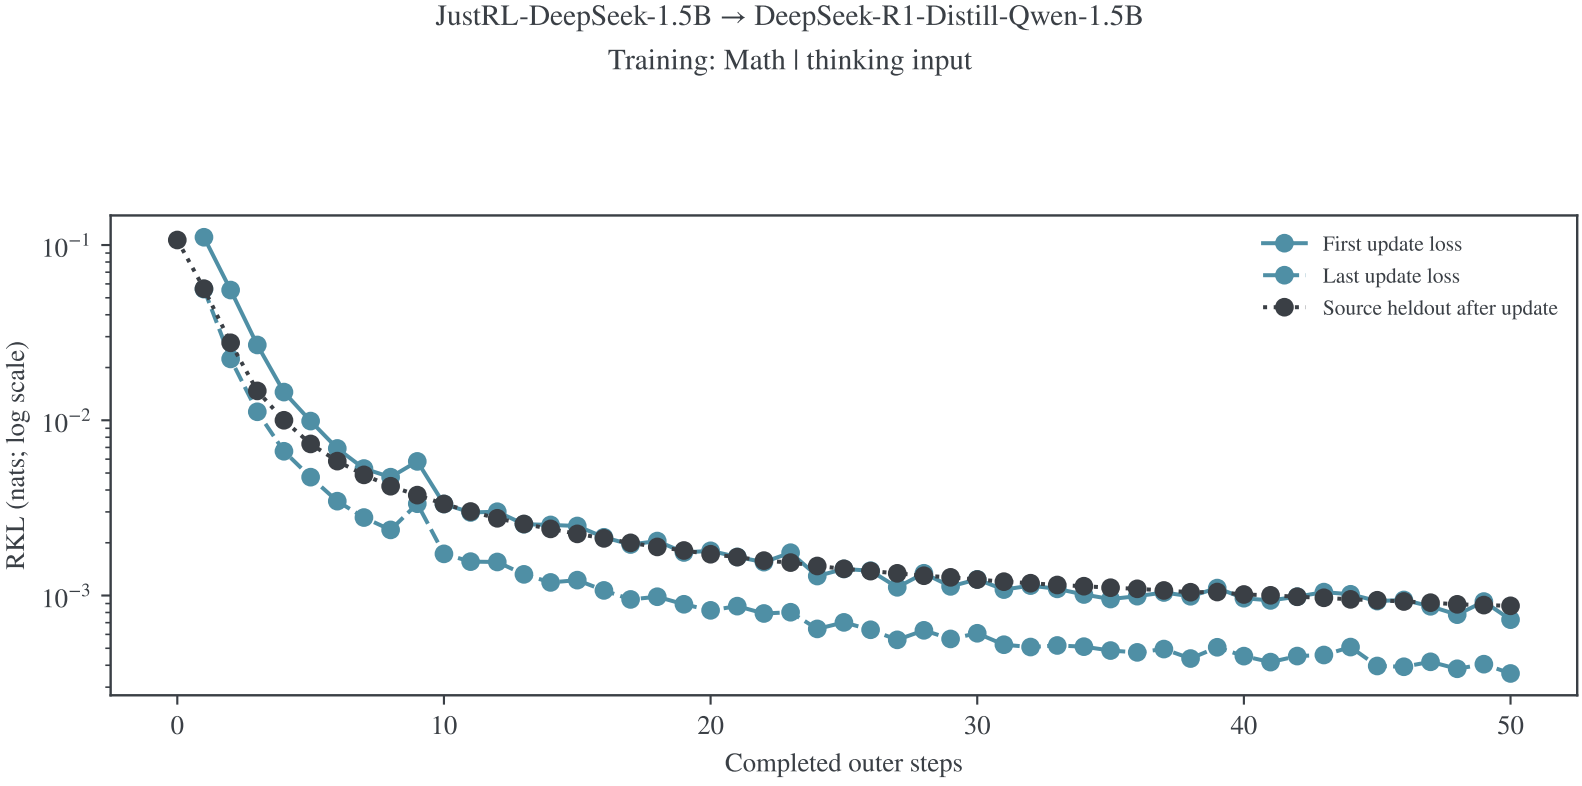
JustRL-DeepSeek-1.5B → DeepSeek-R1-Distill-Qwen-1.5B
Training: Math | thinking input
RKL (nats; log scale)
10^{-1}
10^{-2}
10^{-3}
0
10
20
30
40
50
Completed outer steps
First update loss
Last update loss
Source heldout after update
First/last losses precede their respective updates; loss token-mean, heldout question-mean.