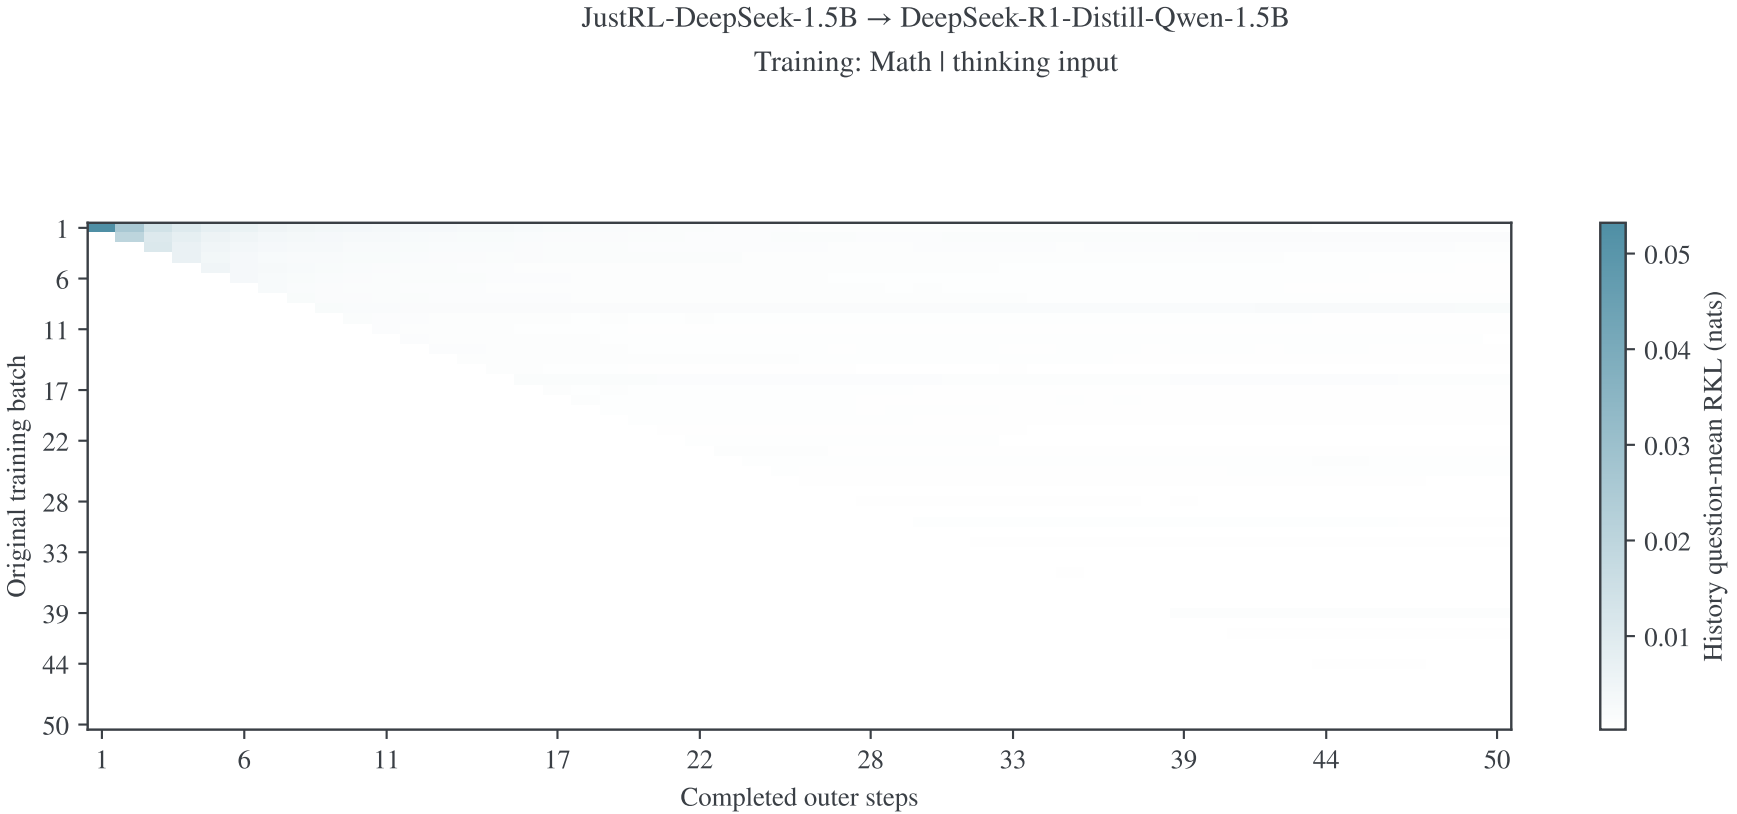
JustRL-DeepSeek-1.5B → DeepSeek-R1-Distill-Qwen-1.5B
Training: Math | thinking input
Original training batch
1
6
11
17
22
28
33
39
44
50
Completed outer steps
History question-mean RKL (nats)
0.05
0.04
0.03
0.02
0.01
Each cell uses the retained random sample from that batch; blank cells are unobserved.