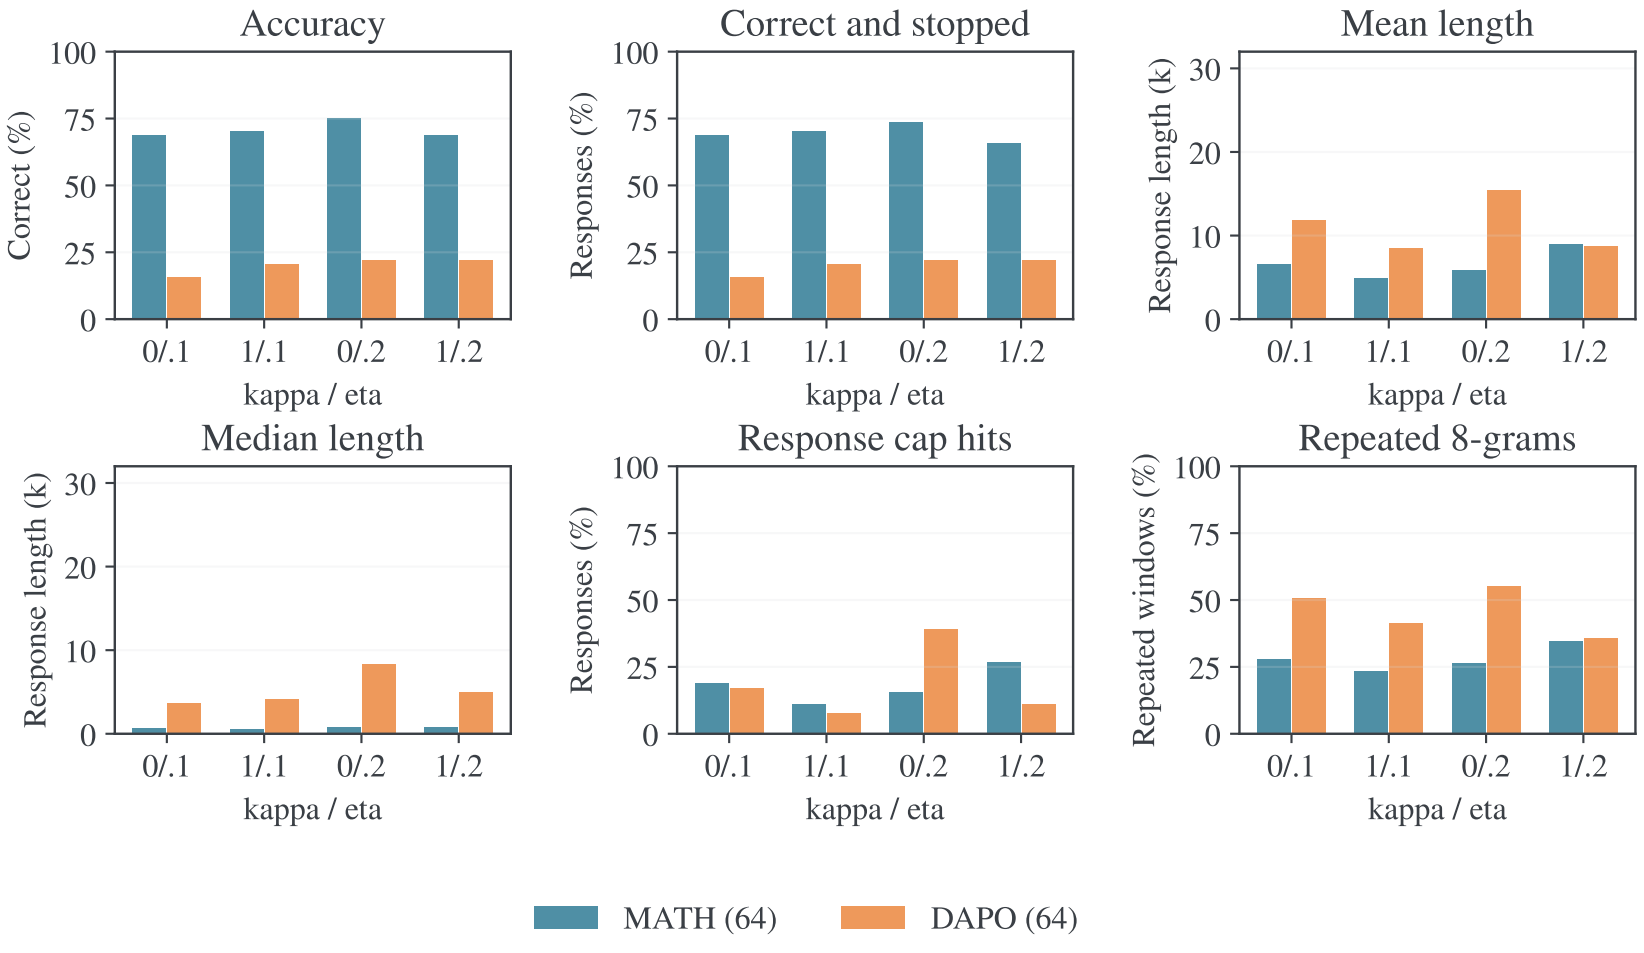
Qwen3-1.7B-Base <- Qwen3-30B-A3B-Instruct-2507
Final-checkpoint evaluation: first four arms
Accuracy
Correct (%)
Correct and stopped
Responses (%)
Mean length
Response length (k)
Median length
Response length (k)
Response cap hits
Responses (%)
Repeated 8-grams
Repeated windows (%)
0/.1
1/.1
0/.2
1/.2
kappa / eta
MATH (64)
DAPO (64)
Outer 5 / 80 updates; 64 questions per dataset, one response each; cap 31,744.
Single-seed pilot; sweep-arm evaluation has not run yet.