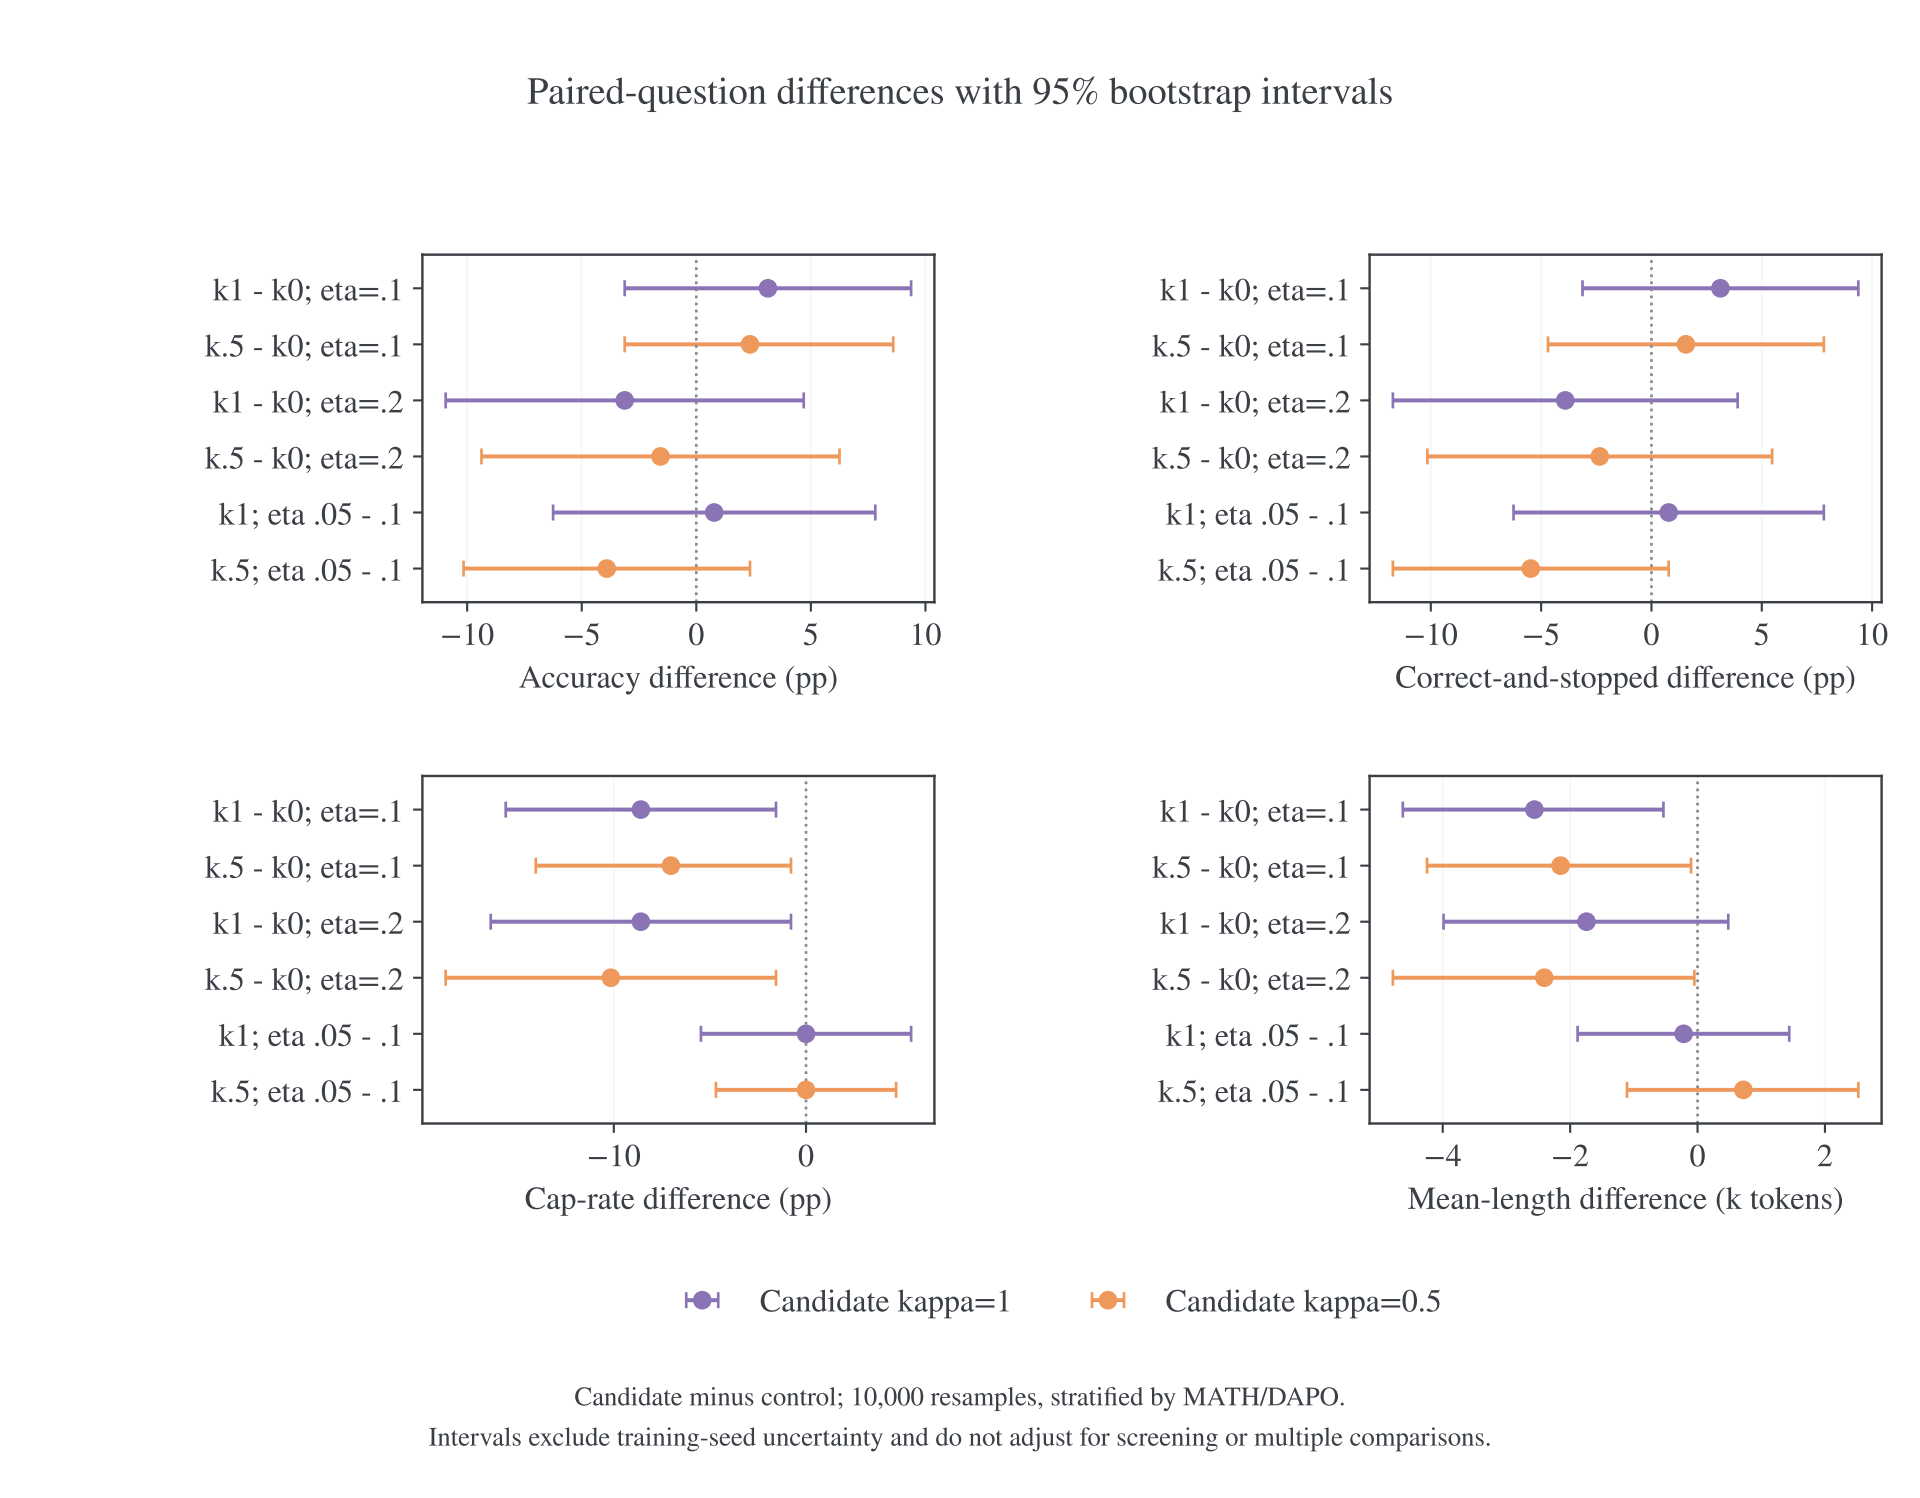
Paired-question differences with 95% bootstrap intervals
k1 - k0; eta=.1
k.5 - k0; eta=.1
k1 - k0; eta=.2
k.5 - k0; eta=.2
k1; eta .05 - .1
k.5; eta .05 - .1
−10
−5
0
5
10
Accuracy difference (pp)
Correct-and-stopped difference (pp)
Cap-rate difference (pp)
−4
−2
2
Mean-length difference (k tokens)
Candidate kappa=1
Candidate kappa=0.5
Candidate minus control; 10,000 resamples, stratified by MATH/DAPO.
Intervals exclude training-seed uncertainty and do not adjust for screening or multiple comparisons.