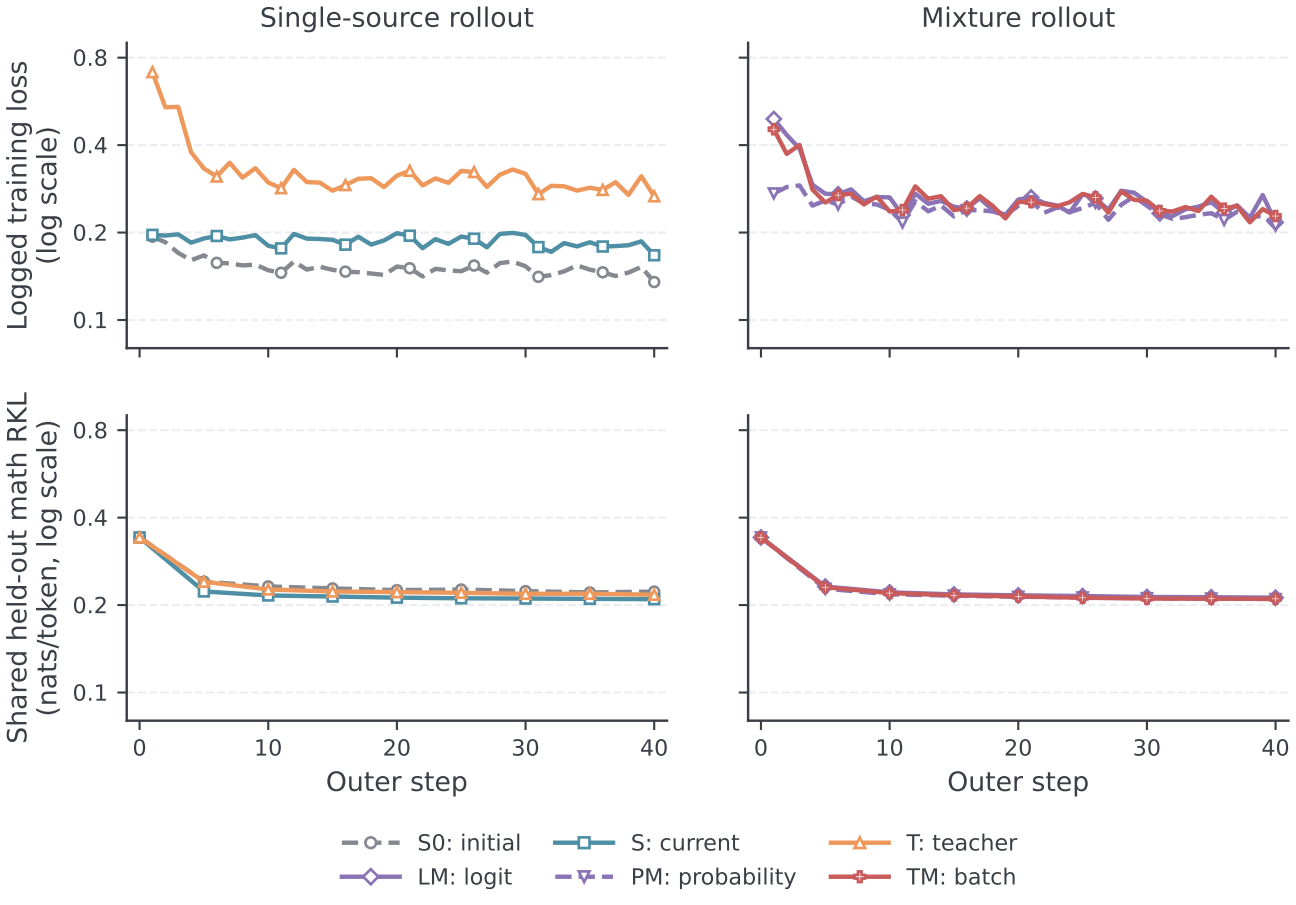
R1-1.5B student | Polaris-7B-Preview teacher
Single-source rollout
Mixture rollout
Logged training loss (log scale)
Shared held-out math RKL (nats/token, log scale)
0.8
0.4
0.2
0.1
0
10
20
30
40
Outer step
S0: initial
S: current
T: teacher
LM: logit
PM: probability
TM: batch
Training: mean of 16 updates. Held-out: common fixed states. No extrapolation.